

Mathematics of Data: From Theory to Computation

Prof. Volkan Cevher
volkan.cevher@epfl.ch

Lecture 10: Adversarial machine learning and generative adversarial networks (GANs)

Laboratory for Information and Inference Systems (LIONS)
École Polytechnique Fédérale de Lausanne (EPFL)

EE-556 (Fall 2021)

lions@epfl

HASLERSTIFTUNG



SDSC



FN-NF



FONDS NATIONAL SUISSE
SCHWEIZERISCHE NATIONALFONDS
FONDO NAZIONALE SVIZZO
SWISS NATIONAL SCIENCE FOUNDATION



erc EPFL

License Information for Mathematics of Data Slides

- ▶ This work is released under a [Creative Commons License](#) with the following terms:
- ▶ **Attribution**
 - ▶ The licensor permits others to copy, distribute, display, and perform the work. In return, licensees must give the original authors credit.
- ▶ **Non-Commercial**
 - ▶ The licensor permits others to copy, distribute, display, and perform the work. In return, licensees may not use the work for commercial purposes – unless they get the licensor's permission.
- ▶ **Share Alike**
 - ▶ The licensor permits others to distribute derivative works only under a license identical to the one that governs the licensor's work.
- ▶ [Full Text of the License](#)

Outline


- ▶ This class
 - ▶ Adversarial Machine Learning (minmax)
 - ▶ Adversarial training
 - ▶ Generative adversarial networks
 - ▶ Difficulty of minmax
- ▶ Next class
 - ▶ Primal-dual optimization (Part 1)

Adversarial machine learning

$$\min_{\mathbf{x} \in \mathcal{X}} \max_{\mathbf{y} \in \mathcal{Y}} \Phi(\mathbf{x}, \mathbf{y})$$

- A seemingly simple optimization formulation
- Critical in machine learning with many applications
 - ▶ Adversarial examples and training
 - ▶ Generative adversarial networks
 - ▶ *Robust reinforcement learning

From empirical risk minimization...



Definition (Empirical Risk Minimization (ERM))

Let $h_{\mathbf{x}} : \mathbb{R}^p \rightarrow \mathbb{R}$ be a model with parameters \mathbf{x} and let $\{(\mathbf{a}_i, b_i)\}_{i=1}^n$ be samples with $b_i \in \{-1, 1\}$ and $\mathbf{a}_i \in \mathbb{R}^p$. The ERM problem reads

$$\min_{\mathbf{x}} \left\{ R_n(x) := \frac{1}{n} \sum_{i=1}^n L(h_{\mathbf{x}}(\mathbf{a}_i), b_i) \right\},$$

where $L(h_{\mathbf{x}}(\mathbf{a}_i), b_i)$ is the loss on the sample (\mathbf{a}_i, b_i) .

Some frequently used loss functions

- ▶ $L(h_{\mathbf{x}}(\mathbf{a}), b) = \log(1 + \exp(-b h_{\mathbf{x}}(\mathbf{a})))$ *logistic loss*
- ▶ $L(h_{\mathbf{x}}(\mathbf{a}), b) = (b - h_{\mathbf{x}}(\mathbf{a}))^2$ *squared error*
- ▶ $L(h_{\mathbf{x}}(\mathbf{a}), b) = \max(0, 1 - b h_{\mathbf{x}}(\mathbf{a}))$ *hinge loss*

...Into adversarial examples

Definition (Adversarial examples [26])

Let $h_{\mathbf{x}^*} : \mathbb{R}^p \rightarrow \mathbb{R}$ be a model trained through empirical risk minimization, with optimal parameters \mathbf{x}^* . Let (\mathbf{a}, b) be a sample with $b \in \{-1, 1\}$ and $\mathbf{a} \in \mathbb{R}^p$. An **adversarial example** is a perturbation $\boldsymbol{\eta} \in \mathbb{R}^n$ designed to lead the trained model $h_{\mathbf{x}^*}$ to misclassify a given input \mathbf{a} . Given an $\epsilon > 0$, it is constructed by solving

$$\boldsymbol{\eta} \in \arg \max_{\boldsymbol{\eta}: \|\boldsymbol{\eta}\| \leq \epsilon} L(h_{\mathbf{x}^*}(\mathbf{a} + \boldsymbol{\eta}), b)$$

Example norms frequently used in adversarial attacks

- ▶ The most commonly used norm is the ℓ_∞ -norm [7, 19].
- ▶ The use of ℓ_1 -norm leads to sparse attacks.



Figure: (Left) An ℓ_∞ -attack: The alteration is hard to perceive. (Right) An ℓ_1 -attack: The alteration in this case is obvious.

Challenge: Robustness to adversarial examples





Figure: Understanding the robustness of a classifier in high-dimensional spaces. Shafahi et al. 2019.

A robustness example: Linear prediction

Linear model

Consider a linear model $h_{\mathbf{x}^*}(\mathbf{a}) = \langle \mathbf{x}^*, \mathbf{a} \rangle$ with weights $\mathbf{x}^* \in \mathbb{R}^p$, for some input \mathbf{a} .

An adversarial perturbation

We aim at finding the perturbation $\boldsymbol{\eta} \in \mathbb{R}^n$ subject to $\|\boldsymbol{\eta}\|_\infty \leq \epsilon$ that produces the largest change on $h_{\mathbf{x}^*}(\mathbf{a})$:

$$\begin{aligned} \max_{\boldsymbol{\eta}: \|\boldsymbol{\eta}\|_\infty \leq \epsilon} h_{\mathbf{x}^*}(\mathbf{a} + \boldsymbol{\eta}) &= \max_{\boldsymbol{\eta}: \|\boldsymbol{\eta}\|_\infty \leq \epsilon} \langle \mathbf{x}^*, \mathbf{a} + \boldsymbol{\eta} \rangle \\ &= \langle \mathbf{x}^*, \mathbf{a} \rangle + \max_{\boldsymbol{\eta}: \|\boldsymbol{\eta}\|_\infty \leq \epsilon} \langle \mathbf{x}^*, \boldsymbol{\eta} \rangle \quad \triangleright \text{As } \mathbf{a} \text{ does not influence the optimization.} \\ &= \langle \mathbf{x}^*, \mathbf{a} \rangle + \max_{\boldsymbol{\eta}: \|\boldsymbol{\eta}\|_\infty \leq 1} \langle \mathbf{x}^*, \epsilon \boldsymbol{\eta} \rangle \quad \triangleright \text{By the change of variables } \boldsymbol{\eta} := \boldsymbol{\eta}/\epsilon \\ &= \langle \mathbf{x}^*, \mathbf{a} \rangle + \epsilon \|\mathbf{x}^*\|_1 \quad \triangleright \text{Definition of the dual norm } \|\mathbf{x}\|_1 := \max_{\boldsymbol{\eta}: \|\boldsymbol{\eta}\|_\infty \leq 1} \langle \mathbf{x}, \boldsymbol{\eta} \rangle \end{aligned}$$

Taking $\boldsymbol{\eta}^* = \text{sign}(\mathbf{x}^*)$ achieves this maximum: $\langle \mathbf{x}, \epsilon \text{ sign}(\mathbf{x}^*) \rangle = \epsilon \sum_{i=1}^n \text{sign}(x_i^*) x_i^* = \epsilon \sum_{i=1}^n |x_i^*| = \epsilon \|\mathbf{x}^*\|_1$.

Remarks:

- For the linear model, we have $\nabla_{\mathbf{a}} h_{\mathbf{x}^*}(\mathbf{a}) = \mathbf{x}^*$.
- *The gradient sign of $h_{\mathbf{x}^*}$ with respect to the input \mathbf{a} achieves the worst perturbation.*
- Sparse models are robust in linear prediction.

Adversarial examples in neural networks

- Target problem:

$$\max_{\eta: \|\eta\|_\infty \leq \epsilon} L(h_{\mathbf{x}^*}(\mathbf{a} + \eta), \mathbf{b})$$

- Historically, researchers first tried to find approximate solutions that empirically perform well [7, 19].

Fast Gradient Sign Method (FGSM) [7]

Let $h_{\mathbf{x}^*} : \mathbb{R}^p \rightarrow \mathbb{R}$ be a model trained through empirical risk minimization on the loss L , with optimal parameters \mathbf{x}^* . Let (\mathbf{a}, b) be a sample with $b \in \{-1, 1\}$ and $\mathbf{a} \in \mathbb{R}^p$. The *Fast Gradient Sign Method* computes the adversarial example

$$\eta = \epsilon \operatorname{sign}(\nabla_{\mathbf{a}} L(h_{\mathbf{x}^*}(\mathbf{a}), b)) = \epsilon \operatorname{sign}(\nabla_{\mathbf{a}} h_{\mathbf{x}^*}(\mathbf{a}) \nabla_h L(h_{\mathbf{x}^*}(\mathbf{a}), b))$$

Remarks:

- The FGSM obtains adversarial examples by using *sign of the gradient of the loss*.
- Such an approach can be viewed as a linearization of the objective L around the data \mathbf{a} .
- For single output $h_{\mathbf{x}}(\mathbf{a})$, $\nabla_h L(h_{\mathbf{x}^*}(\mathbf{a}), b)$ is a scalar,
 - ▶ $\operatorname{sign}(\nabla_{\mathbf{a}} h_{\mathbf{x}^*}(\mathbf{a}))$ pattern is important

Results of FGSM on MNIST



Figure: MNIST images with the predicted digit.

Figure: MNIST images perturbed by a FGSM attack.

Taken from https://adversarial-ml-tutorial.org/adversarial_examples/

Adversarial examples and proximal gradient descent

- Target problem:

$$\max_{\boldsymbol{\eta}: \|\boldsymbol{\eta}\|_\infty \leq \epsilon} L(h_{\mathbf{x}^*}(\mathbf{a} + \boldsymbol{\eta}), \mathbf{b})$$

- We can do better than FGSM via proximal gradient methods for composite minimization:

$$\max_{\boldsymbol{\eta} \in \mathbb{R}^p} \underbrace{L(h_{\mathbf{x}^*}(\mathbf{a} + \boldsymbol{\eta}), \mathbf{b})}_{f(\boldsymbol{\eta})} + \underbrace{\delta_{\mathcal{N}}(\boldsymbol{\eta})}_{g(\boldsymbol{\eta})},$$

where $\delta_{\mathcal{N}}(\boldsymbol{\eta})$ is the indicator function of the ball $\mathcal{N} := \{\boldsymbol{\eta} : \|\boldsymbol{\eta}\|_\infty \leq \epsilon\}$.

Recall: Proximal operator of indicator functions

For the indicator functions of simple sets, e.g., $g(\boldsymbol{\eta}) := \delta_{\mathcal{N}}(\boldsymbol{\eta})$, the prox-operator is the projection operator

$$\text{prox}_{\lambda g}(\boldsymbol{\eta}) := \pi_{\mathcal{N}}(\boldsymbol{\eta}),$$

where $\pi_{\mathcal{N}}(\boldsymbol{\eta})$ denotes the projection of $\boldsymbol{\eta}$ onto \mathcal{N} . When $\mathcal{N} = \{\boldsymbol{\eta} : \|\boldsymbol{\eta}\|_\infty \leq \lambda\}$, $\pi_{\mathcal{N}}(\boldsymbol{\eta}) = \text{clip}(\boldsymbol{\eta}, [-\lambda, \lambda])$.

Adversarial examples and proximal gradient descent (cont'd)

- Target non-convex problem:

$$\max_{\boldsymbol{\eta} \in \mathbb{R}^p} \underbrace{L(h_{\mathbf{x}^*}(\mathbf{a} + \boldsymbol{\eta}), \mathbf{b})}_{f(\boldsymbol{\eta})} + \underbrace{\delta_{\mathcal{N}}(\boldsymbol{\eta})}_{g(\boldsymbol{\eta})},$$

where $\delta_{\mathcal{N}}(\boldsymbol{\eta})$ is the indicator function of the ball $\mathcal{N} := \{\mathbf{y} : \|\mathbf{y}\|_\infty \leq \epsilon\}$.

Proximal gradient ascent (PGA)

1. Choose $\boldsymbol{\eta}^0 \in \text{dom } f(\boldsymbol{\eta}) + g(\boldsymbol{\eta})$ as initialization.
2. For $k = 0, 1, \dots$, generate a sequence $\{\boldsymbol{\eta}^k\}_{k \geq 0}$ as:

$$\boldsymbol{\eta}^{k+1} := \text{prox}_{\alpha_k g} \left(\boldsymbol{\eta}^k + \alpha_k \nabla f(\boldsymbol{\eta}^k) \right).$$

Remarks:

- PGA results in more powerful adversarial “attacks” than FGSM [13].
- The PGA is incorrectly referred to as projected gradient descent in this literature.
- Practitioners prefer to use several steps of FGSM instead of PGA [15, 16, 19]:

$$\boldsymbol{\eta}^{k+1} = \pi_{\mathcal{X}} \left(\boldsymbol{\eta}^k + \alpha_k \text{sign} (\nabla f(\boldsymbol{\eta}^k)) \right).$$

A proposed link between FGSM and PGD

o Recall

- ▶ A single step of PGA reads $\eta_{\text{PGA}}^{k+1} := \pi_{\mathcal{N}}(\eta^k + \alpha \nabla f(\eta))$
- ▶ The FGSM attack is defined as $\eta_{\text{FGSM}} := \epsilon \text{ sign}(\nabla_{\mathbf{a}} L(h_{\mathbf{x}^*}(\mathbf{a}), \mathbf{b}))$
- ▶ When $\mathcal{N} = \{\eta : \|\eta\|_\infty \leq \lambda\}$, $\pi_{\mathcal{N}}(\eta) = \text{clip}(\eta, [-\lambda, \lambda])$

FGSM as one step of PGA

Let $\eta^0 = \mathbf{0}$ and $\alpha > 0$ such that $(\alpha |\nabla f(\mathbf{0})|)_i > \epsilon$ for $i = 1, \dots, n$. Then, one step of PGA yields

$$\begin{aligned}\eta_{\text{PGD}}^1 &= \pi_{\mathcal{N}}(\eta^0 + \alpha \nabla_{\eta} \nabla f(\eta^0)) \\ &= \text{clip}(\alpha \nabla f(\mathbf{0}), [-\epsilon, \epsilon]) \quad \triangleright \eta^0 = \mathbf{0} \\ &= \epsilon \text{ sign}(\nabla f(\mathbf{0})) \quad \triangleright \text{All values are outside of the interval } [-\epsilon, \epsilon] \\ &= \epsilon \text{ sign}(\nabla_{\mathbf{a}} L(h_{\mathbf{x}^*}(\mathbf{a}), \mathbf{b})) = \eta_{\text{FGSM}} \quad \triangleright \nabla f(\mathbf{0}) = \nabla_{\mathbf{a}} L(h_{\mathbf{x}^*}(\mathbf{a}), \mathbf{b})\end{aligned}$$

A proposed link between FGSM and PGD

o Recall

- ▶ A single step of PGA reads $\eta_{\text{PGA}}^{k+1} := \pi_{\mathcal{N}}(\eta^k + \alpha \nabla f(\eta))$
- ▶ The FGSM attack is defined as $\eta_{\text{FGSM}} := \epsilon \text{ sign}(\nabla_{\mathbf{a}} L(h_{\mathbf{x}^*}(\mathbf{a}), \mathbf{b}))$
- ▶ When $\mathcal{N} = \{\eta : \|\eta\|_\infty \leq \lambda\}$, $\pi_{\mathcal{N}}(\eta) = \text{clip}(\eta, [-\lambda, \lambda])$



FGSM as one step of PGA

Let $\eta^0 = \mathbf{0}$ and $\alpha > 0$ such that $(\alpha |\nabla f(\mathbf{0})|)_i > \epsilon$ for $i = 1, \dots, n$. Then, one step of PGA yields

$$\begin{aligned}\eta_{\text{PGD}}^1 &= \pi_{\mathcal{N}}(\eta^0 + \alpha \nabla_{\eta} \nabla f(\eta^0)) \\ &= \text{clip}(\alpha \nabla f(\mathbf{0}), [-\epsilon, \epsilon]) \quad \triangleright \eta^0 = \mathbf{0} \\ &= \epsilon \text{ sign}(\nabla f(\mathbf{0})) \quad \triangleright \text{All values are outside of the interval } [-\epsilon, \epsilon] \\ &= \epsilon \text{ sign}(\nabla_{\mathbf{a}} L(h_{\mathbf{x}^*}(\mathbf{a}), \mathbf{b})) = \eta_{\text{FGSM}} \quad \triangleright \nabla f(\mathbf{0}) = \nabla_{\mathbf{a}} L(h_{\mathbf{x}^*}(\mathbf{a}), \mathbf{b})\end{aligned}$$

Multiple steps of FGSM: A connection to majorization-minimization in Lecture 3

Minimization-majorization for concave functions

Let f be a concave function which is smooth in the ℓ_∞ -norm with constant L_∞ . Our target non-convex problem is given by

$$\max_{\eta} f(\eta) + \delta_{\mathcal{N}}(\eta)$$

where $\delta_{\mathcal{N}}(\eta)$ is the indicator function of the ball $\mathcal{N} := \{\eta : \|\eta\|_\infty \leq \epsilon\}$. Smoothness in ℓ_∞ -norm implies

$$f(\eta) + \delta_{\mathcal{N}}(\eta) \geq f(\zeta) + \underbrace{\langle \nabla_{\eta} f(\zeta), \eta - \zeta \rangle - \frac{L_\infty}{2} \|\eta - \zeta\|_\infty^2}_{\eta^* \leftarrow \arg \max_{\eta}} + \delta_{\mathcal{X}}(\eta).$$

Maximizing the RHS with respect to η leads to the following (non trivial) solution [4]:

$$\eta^* = \text{clip}(\zeta - t^* \text{sign}(\nabla f(\zeta)), [-\epsilon, \epsilon])$$

where $t^* = \arg \max_{t: \|\eta - \zeta\|_\infty \leq t} \max_{\zeta: \|\zeta\|_\infty \leq \epsilon} \langle \nabla f(\zeta), \eta - \zeta \rangle$ can be found by linear search.

Remarks:

- Setting $\zeta = \eta^k$ and $\eta^* = \eta^{k+1}$ with a fixed step size $\alpha = t^*$, we obtain the update in [15, 16, 19]
$$\eta^{k+1} = \text{clip}\left(\eta^k - t^* \text{sign}(\nabla f(\eta^k)), [-\epsilon, \epsilon]\right).$$

- This proof holds for **concave** and smooth functions, and need further quantification for our setting.

Towards adversarial training

Adversarial Training

Let $h_{\mathbf{x}} : \mathbb{R}^n \rightarrow \mathbb{R}$ be a model with parameters \mathbf{x} and let $\{(\mathbf{a}_i, \mathbf{b}_i)\}_{i=1}^n$, with the data $\mathbf{a}_i \in \mathbb{R}^p$ and the labels \mathbf{b}_i . The problem of adversarial training is the following adversarial optimization problem

$$\min_{\mathbf{x}} \frac{1}{n} \sum_{i=1}^n \left[\max_{\boldsymbol{\eta}: \|\boldsymbol{\eta}\|_\infty \leq \epsilon} L(h_{\mathbf{x}}(\mathbf{a}_i + \boldsymbol{\eta}), \mathbf{b}_i) \right] \approx \min_{\mathbf{x}} \mathbb{E}_{(\mathbf{a}, \mathbf{b}) \sim \mathbb{P}} \left[\max_{\boldsymbol{\eta}: \|\boldsymbol{\eta}\|_\infty \leq \epsilon} L(h_{\mathbf{x}}(\mathbf{a}_i + \boldsymbol{\eta}), \mathbf{b}_i) \right].$$

Note the similarity with the template $\min_{\mathbf{x} \in \mathcal{X}} \max_{\mathbf{y} \in \mathcal{Y}} \Phi(\mathbf{x}, \mathbf{y})$.

Solving the outer problem

Adversarial Training

Let $h_{\mathbf{x}} : \mathbb{R}^n \rightarrow \mathbb{R}$ be a model with parameters \mathbf{x} and let $\{(\mathbf{a}_i, \mathbf{b}_i)\}_{i=1}^n$, with $\mathbf{a}_i \in \mathbb{R}^p$ and \mathbf{b}_i be the corresponding labels. The adversarial training optimization problem is given by

$$\min_{\mathbf{x}} \left\{ \frac{1}{n} \sum_{i=1}^n f_i(\mathbf{x}) := \frac{1}{n} \sum_{i=1}^n \underbrace{\left[\max_{\boldsymbol{\eta}: \|\boldsymbol{\eta}\|_\infty \leq \epsilon} L(h_{\mathbf{x}}(\mathbf{a}_i + \boldsymbol{\eta}), \mathbf{b}_i) \right]}_{=: f_i(\mathbf{x})} \right\}.$$

Note that L is not continuously differentiable due to ReLU, max-pooling, etc.

Solving the outer problem

Adversarial Training

Let $h_{\mathbf{x}} : \mathbb{R}^n \rightarrow \mathbb{R}$ be a model with parameters \mathbf{x} and let $\{(\mathbf{a}_i, \mathbf{b}_i)\}_{i=1}^n$, with $\mathbf{a}_i \in \mathbb{R}^p$ and \mathbf{b}_i be the corresponding labels. The adversarial training optimization problem is given by

$$\min_{\mathbf{x}} \left\{ \frac{1}{n} \sum_{i=1}^n f_i(\mathbf{x}) := \frac{1}{n} \sum_{i=1}^n \underbrace{\left[\max_{\boldsymbol{\eta}: \|\boldsymbol{\eta}\|_\infty \leq \epsilon} L(h_{\mathbf{x}}(\mathbf{a}_i + \boldsymbol{\eta}), \mathbf{b}_i) \right]}_{=: f_i(\mathbf{x})} \right\}.$$

Note that L is not continuously differentiable due to ReLU, max-pooling, etc.

Question

How can we compute the gradient

$$\nabla_{\mathbf{x}} f_i(\mathbf{x}) := \nabla_{\mathbf{x}} \left(\max_{\boldsymbol{\eta}: \|\boldsymbol{\eta}\|_\infty \leq \epsilon} L(h_{\mathbf{x}}(\mathbf{a}_i + \boldsymbol{\eta}), \mathbf{b}_i) \right)?$$

- **Challenge:** It involves differentiating with respect to a maximization.
- **A solution:** We can use Danskin's theorem under some conditions.

Danskin's theorem

Danskin's theorem (Bertsekas variant)

Let $\Phi(\mathbf{x}, \mathbf{y}) : \mathbb{R}^p \times \mathcal{Y} \rightarrow \mathbb{R}$, where $\mathcal{Y} \subset \mathbb{R}^m$ is a compact set and define $f(\mathbf{x}) := \max_{\mathbf{y} \in \mathcal{Y}} \Phi(\mathbf{x}, \mathbf{y})$. Suppose that $\Phi(\mathbf{x}, \mathbf{y})$ is convex for each \mathbf{y} in the compact set \mathcal{Y} ; the interior of the domain of f is nonempty; and $\Phi(\mathbf{x}, \mathbf{y})$ is continuous.

Define $\mathcal{Y}^*(\mathbf{x}) := \arg \max_{\mathbf{y} \in \mathcal{Y}} \Phi(\mathbf{x}, \mathbf{y})$ as the set of maximizers and $\mathbf{y}^* \in \mathcal{Y}^*$ as an element of this set. We have

1. $f(\mathbf{x})$ is a convex function.
2. If $\mathcal{Y}^*(\mathbf{x})$ is a singleton, then the function $f(\mathbf{x}) = \max_{\mathbf{y} \in \mathcal{Y}} \Phi(\mathbf{x}, \mathbf{y})$ is differentiable at \mathbf{x} :

$$\nabla_{\mathbf{x}} f(\mathbf{x}) = \nabla_{\mathbf{x}} \left(\max_{\mathbf{y} \in \mathcal{Y}} \phi(\mathbf{x}, \mathbf{y}) \right) = \nabla_{\mathbf{x}} \Phi(\mathbf{x}, \mathbf{y}^*).$$

3. If $\mathcal{Y}^*(\mathbf{x})$ contains more than one element, then the subdifferential $\partial_{\mathbf{x}} f(\mathbf{x})$ of f is given by

$$\partial_{\mathbf{x}} f(\mathbf{x}) = \text{conv} \{ \partial_{\mathbf{x}} \Phi(\mathbf{x}, \mathbf{y}^*) : \mathbf{y}^* \in \mathcal{Y}^*(\mathbf{x}) \}.$$

Remarks:

- o The adversarial problem is not convex in \mathbf{x} in general.
- o (Sub)Gradients of f are calculated as $\nabla_{\mathbf{x}} f(\mathbf{x}) = \nabla_{\mathbf{x}} \Phi(\mathbf{x}, \mathbf{y}^*)$.

The adversarial training formulation

Adversarial Training

Let $h_{\mathbf{x}} : \mathbb{R}^n \rightarrow \mathbb{R}$ be a model with parameters \mathbf{x} and let $\{(\mathbf{a}_i, \mathbf{b}_i)\}_{i=1}^n$, with $\mathbf{a}_i \in \mathbb{R}^p$ and \mathbf{b}_i be the corresponding labels. The adversarial training optimization problem is given by

$$\min_{\mathbf{x}} \left\{ \frac{1}{n} \sum_{i=1}^n f_i(\mathbf{x}) := \frac{1}{n} \sum_{i=1}^n \underbrace{\left[\max_{\eta: \|\eta\|_\infty \leq \epsilon} L(h_{\mathbf{x}}(\mathbf{a}_i + \eta), \mathbf{b}_i) \right]}_{=: f_i(\mathbf{x})} \right\}.$$

L is not differentiable due to non-smooth activation functions (ReLU), nor convex in \mathbf{x} because of the neural network structure.




Figure: Descent directions in 2D should be an element of the cone of descent directions $D(f, \cdot)$.

Descent Directions in the non-convex case

General Danskin's Theorem

Assume \mathcal{Y} is compact and $\Phi(\mathbf{x}, \mathbf{y})$ differentiable in \mathbf{x} but not necessarily convex in \mathbf{x} . Define $\mathcal{Y}^*(\mathbf{x}) := \arg \max_{\mathbf{y} \in \mathcal{Y}} \Phi(\mathbf{x}, \mathbf{y})$ as the set of maximizers. Then $f(\mathbf{x}) := \max_{\mathbf{y} \in \mathcal{Y}} \Phi(\mathbf{x}, \mathbf{y})$ is *directionally differentiable* and its directional derivative is given by

$$Df(\mathbf{x}, \mathbf{d}) = \max_{\mathbf{y}^* \in \mathcal{Y}^*(\mathbf{x})} \langle \mathbf{d}, \nabla_{\mathbf{x}} \Phi(\mathbf{x}, \mathbf{y}^*) \rangle \quad (1)$$

Corollary (Corollary A.2 in [19])

Let \mathbf{y}_0^* be a maximizer of $\max_{\mathbf{y} \in \mathcal{Y}} \Phi(\mathbf{x}, \mathbf{y})$. Then as long as $\nabla_{\mathbf{x}} \Phi(\mathbf{x}, \mathbf{y}_0^*)$ is non-zero, $-\nabla_{\mathbf{x}} \Phi(\mathbf{x}, \mathbf{y}_0^*)$ is a descent direction for $f(\mathbf{x})$.

Caveat

What is the definition of (i) directional derivative and (ii) descent direction?

A practical implementation of adversarial training: Stochastic subgradient descent

Stochastic Adversarial Training [19]	
Input: learning rate α_k , iterations T , batch size K .	
1.	initialize neural network parameters \mathbf{x}^0
2. For $k = 0, 1, \dots, T$:	
i.	initialize update vector $\mathbf{g}^k := 0$
ii.	select a mini-batch of data $B \subset \{1, \dots, n\}$ with $ B = K$
iii. For $i \in B$:	
a.	Find an attack $\boldsymbol{\eta}^*$ by (approximately) solving $\boldsymbol{\eta}^* \in \arg \max_{\boldsymbol{\eta}: \ \boldsymbol{\eta}\ _\infty \leq \epsilon} L(h_{\mathbf{x}^k}(\mathbf{a}_i + \boldsymbol{\eta}), \mathbf{b}_i)$
b.	Store update $\mathbf{g}^k := \mathbf{g}^k + \nabla_{\mathbf{x}} L(h_{\mathbf{x}^k}(\mathbf{a}_i + \boldsymbol{\eta}^*), \mathbf{b}_i)$
iv.	Update parameters $\mathbf{x}^{k+1} := \mathbf{x}^k - \frac{\alpha_k}{K} \mathbf{g}^k$

Remarks:

- Expensive but worth it!
- Inner problem iii.a cannot be solved to optimality (non-convex).
- Practitioners use FGSM or PGA or PGA- ℓ_∞ to approximate the true $\boldsymbol{\eta}^*$.
- Update in step iii.b is motivated by Corollary A.2 in [19]

Application: Adversarial training for better interpretability

- Retinopathy classification problem: Given a retinal image (left), predict whether there is a disease.
- **Zeiss:** How can we interpret the prediction of a model $h_x(\mathbf{a})$?
- **Solution:** Look at $\nabla_{\mathbf{x}} h_x(\mathbf{a})$, called the saliency map [25]. Adversarial training helps!







Table: **Left:** Ground truth image, **Middle:** Saliency map, **Right:** Saliency map with adversarial training.

Adversarial machine learning: Introduction to Generative Adversarial Networks (GANs)

- Recall the parametric density estimation setting




(source: <http://mmlab.ie.cuhk.edu.hk/projects/CelebA.html>)

$$\mathbf{a}_i = [\dots \text{images} \dots]$$

$$b_i = [\dots \text{probability} \dots]$$

- Goal: Games, denoising, image recovery...



- Generator $\mathbb{P}_{\mathbf{a}}$
 - Nature
- Supervisor $\mathbb{P}_{B|\mathbf{a}}$
 - Frequency data
- Learning Machine $h_{\mathbf{x}}(\mathbf{a}_i)$
 - Data scientist: Mathematics of Data

A notion of distance between distributions




Figure: The Earth Mover's distance

Minimum cost transportation problem (Monge's problem)

Find a *transport map* $T : \mathbb{R}^d \rightarrow \mathbb{R}^d$ such that $T(X) \sim Y$, minimizing the cost

$$\text{cost}(T) := \mathbf{E}_X \|X - T(X)\|. \quad (2)$$

The Wasserstein distance

Definition

Let μ and ν be two probability measures on \mathbb{R}^d . Their set of couplings is defined as

$$\Gamma(\mu, \nu) := \{\pi \text{ prob. measure on } \mathbb{R}^d \times \mathbb{R}^d \text{ with marginals } \mu, \nu\} \quad (3)$$

Definition (q -Wasserstein distance (Primal))

$$W_q(\mu, \nu) := \left(\inf_{\pi \in \Gamma(\mu, \nu)} \mathbf{E}_{(\mathbf{a}, \mathbf{a}') \sim \pi} d(\mathbf{a}, \mathbf{a}')^q \right)^{1/q} \quad (4)$$

where $q = 1, 2$ and d is a distance.




Figure: Two one-dimensional distributions plotted on the x and y axes, and one possible joint distribution that defines a transport plan between them (https://en.wikipedia.org/wiki/Wasserstein_metric).

Properties of the Wasserstein distance

- For any $q \geq 1$, the q -Wasserstein distance *is* a distance:
 - ▶ $W_q(\mu, \nu) = 0$ if and only if μ, ν have the same density almost everywhere (identity).
 - ▶ $W_q(\mu, \nu) = W_q(\nu, \mu)$ (symmetry).
 - ▶ $W_q(\mu, \rho) \leq W_q(\mu, \nu) + W_q(\nu, \rho)$ (triangle inequality).

Problem (Wasserstein Projection)

Given a target probability measure μ on \mathbb{R}^d we are interested in solving the following optimization problem:

$$\min_{\nu \in \Delta} W_q(\mu, \nu), \quad (5)$$

where Δ is a set of probability measures on \mathbb{R}^d , and q is often selected as 1 or 2.

A way to model complex distributions: The push-forward measure

- Traditionally, we use analytical distributions: Restricts what we could model in real applications.
- Now, we use more expressive probability measures via *push-forward measures* with neural networks

Definition

- Let $\omega \sim p_\Omega$ be a random variable.
- $h_x(\cdot) : \mathbb{R}^p \rightarrow \mathbb{R}^m$ a function parameterized by parameters x .

The pushforward measure of p_Ω under h_x , denoted by $h_x \# p_\Omega$ is the distribution of $h_x(\omega)$.

Example: Chi-square distribution

Let $\omega \sim p_\Omega := \mathcal{N}(0, 1)$ be the normal distribution. Let $h_x : \mathbb{R} \rightarrow \mathbb{R}$, $h_x(\omega) = w^x$. Let us fix $x = 2$. Then, $h_x \# p_\Omega$ is the chi-square distribution with one degree of freedom.

Explanation: Change of variables.

Assume that $h : \mathbb{R}^n \rightarrow \mathbb{R}^n$ is monotonic. Given the random variable $\omega \sim p_\Omega$ with probability density function $p_\Omega(\omega)$, the density $p_Y(y)$ of $y = h_x(\omega)$ reads

$$p_Y(y) = p_\Omega(h_x^{-1}(y)) \det(\mathbf{J}_y h_x^{-1}(y))$$

where \det denotes the determinant operation.

Towards an optimization problem

Problem (Ideal parametric density estimator)

Given a true distribution μ^\natural , we can solve the following optimization problem,

$$\min_{\mathbf{x}} W_1(\mu^\natural, h_{\mathbf{x}} \# p_\Omega), \quad (6)$$

where the measurable function $h_{\mathbf{x}}$ is parameterized by \mathbf{x} and $\omega \sim p_\Omega$ is “simple” e.g., Gaussian.

- Issues:

- We only have access to empirical samples $\hat{\mu}_n$ of μ^\natural .
- W_1 is non-smooth, it cannot be computed exactly.





Figure: Schematic of a generative model, $h_{\mathbf{x}} \# \omega$ [6, 12].

Learning without concentration

- We can minimize $W_1(\hat{\mu}_n, h_{\mathbf{x}} \# p_{\Omega})$ with respect to \mathbf{x} .
- Figure: Empirical distribution (blue), $\hat{\mu}_n = \sum_{i=1}^n \delta_i$



A plug-in empirical estimator

Using the triangle inequality for Wasserstein distances we can upper bound in the follow way,

$$W_1(\mu^\natural, h_{\mathbf{x}} \# p_{\Omega}) \leq W_1(\mu^\natural, \hat{\mu}_n) + W_1(\hat{\mu}_n, h_{\mathbf{x}} \# p_{\Omega}), \quad (7)$$

where $\hat{\mu}_n$ is the empirical estimator of μ^\natural obtained from n independent samples from μ^\natural .

Theorem (Slow convergence of empirical measures in 1-Wasserstein [27, 3])

Let μ^\natural be a measure defined on \mathbb{R}^p and let $\hat{\mu}_n$ be its empirical measure. Then the $\hat{\mu}_n$ converges, in the worst case, at the following rate,

$$W_1(\mu^\natural, \hat{\mu}_n) \gtrsim n^{-1/p}. \quad (8)$$

Remarks:

- Using an empirical estimator in high-dimensions is terrible in the worst case.
- However, it does not directly say that $W_1(\mu^\natural, h_{\mathbf{x}} \# p_{\Omega})$ will be large.
- So we can still proceed and hope our parameterization interpolates harmlessly.

Duality of 1-Wasserstein

- Instead of computing W_1 , we can obtain lower bounds using duality.

Theorem (Kantorovich-Rubinstein duality)

$$W_1(\mu, \nu) = \sup_{\mathbf{d}} \{ \langle \mathbf{d}, \mu \rangle - \langle \mathbf{d}, \nu \rangle : \mathbf{d} \text{ is 1-Lipschitz} \} \quad (9)$$

Remark: \mathbf{d} is the “dual” variable. In the literature, it is commonly referred to as the “discriminator.”

Inner product is an expectation

$$\langle \mathbf{d}, \mu \rangle = \int \mathbf{d} d\mu = \int \mathbf{d}(\mathbf{a}) d\mu(\mathbf{a}) = E_{\mathbf{a} \sim \mu} [\mathbf{d}(\mathbf{a})]. \quad (10)$$

Kantorovich-Rubinstein duality applied to our objective

$$W_1(\hat{\mu}_n, h_{\mathbf{x}} \# \omega) = \sup \left\{ E_{\mathbf{a} \sim \hat{\mu}_n} [\mathbf{d}(\mathbf{a})] - E_{\mathbf{a} \sim h_{\mathbf{x}} \# \omega} [\mathbf{d}(\mathbf{a})] : \mathbf{d} \text{ is 1-Lipschitz} \right\} \quad (11)$$

Integral Probability Metrics

We can define a more general class of (semi)metrics in the space of probability distributions

Definition (Integral Probability Metric)

Let \mathcal{F} be a class of functions from \mathbb{R}^p to \mathbb{R} . For two probability measures μ and ν , the IPM associated to \mathcal{F} is defined as:

$$\mathcal{F}(\mu, \nu) := \sup_{f \in \mathcal{F}} \langle f, \mu \rangle - \langle f, \nu \rangle = \sup_{f \in \mathcal{F}} E_{\mathbf{a} \sim \mu}[f(\mathbf{a})] - E_{\mathbf{a} \sim \nu}[f(\mathbf{a})] \quad (12)$$

- Remarks:**
- The 1-Wasserstein distance corresponds to $\mathcal{F} := \{f : \mathbb{R}^p \rightarrow \mathbb{R}, f \text{ is } 1-\text{Lipschitz}\}$
 - The class cannot be described with finite parameters.

Neural network distances inspired by the 1-Wasserstein distance

- We use neural networks to parametrize a class of functions.
- Constraining the Lipschitz constant of Neural Networks is NP-Hard [22].
- We can constrain upper bounds on the Lipschitz constant [17].

Lemma

Let $h_{\mathbf{X}_1, \mathbf{X}_2}(\mathbf{a}) := \mathbf{X}_2^T \sigma(\mathbf{X}_1 \mathbf{a})$ be a one-hidden-layer neural network. Then its Lipschitz constant $L_{\mathbf{X}_1, \mathbf{X}_2}$ with respect to the ℓ_2 -norm is bounded as:

$$L_{\mathbf{X}_1, \mathbf{X}_2} \leq \|\mathbf{X}_1\|_2 \|\mathbf{X}_2\|_2 \quad (13)$$

Neural Network Distance

Let

$$\mathcal{F} := \{h_{\mathbf{X}_1, \mathbf{X}_2}(\mathbf{a}) = \mathbf{X}_2^T \sigma(\mathbf{X}_1 \mathbf{a}) : \|\mathbf{X}_2\|_2 \leq 1, \|\mathbf{X}_1\|_2 \leq 1\}. \quad (14)$$

The IPM corresponding to \mathcal{F} is referred to as a *Neural Network Distance*.

Remark: ◦ Different network architectures/constraints lead to different Neural Network distance notions.

Wasserstein GANs formulation

- Ingredients:

- ▶ fixed *noise* distribution p_Ω (e.g., normal)
- ▶ target distribution $\hat{\mu}_n$ (natural images)
- ▶ \mathcal{X} parameter class inducing a class of functions (generators)
- ▶ \mathcal{Y} parameter class inducing a class of functions (dual variables)

Wasserstein GANs formulation [1]

Define a parameterized function $d_y(a)$, where $y \in \mathcal{Y}$ such that $d_y(a)$ is 1-Lipschitz. In this case, the Wasserstein GAN optimization problem is given by

$$\min_{x \in \mathcal{X}} \left(\max_{y \in \mathcal{Y}} E_{a \sim \hat{\mu}_n} [d_y(a)] - E_{\omega \sim p_\Omega} [d_y(h_x(\omega))] \right). \quad (15)$$

General diagram of GANs




Figure: Generator/dual variable/dataset relation in GANs

The theory-practice gap: Enforcing 1-Lipschitz of the discriminator

Weight clipping [1]


The “dual” or the “discriminator” d_y weights y are constrained by an ℓ_∞ -ball with radius $c > 0$, denoted as \mathcal{B} , at every iteration with

$$\pi_{\mathcal{B}}(y) = \text{clip}(y, [-c, c]). \quad (16)$$

This trick is used to pseudo-enforce the constraint.

Remark:

- “Weight clipping is a clearly terrible way to enforce a Lipschitz constraint” – original authors.



Gradient penalty [8]

Recall that 1-Lipschitz is equivalent to $\|\nabla_a d_y(a)\|_* \leq 1$. This can be enforced directly through

$$E_{a \sim \hat{\mu}_n} [d_y(a)] - E_{\omega \sim \Omega} [d_y(h_x(\omega))] + \lambda E_{a \sim \nu} [(\|\nabla_a d_y(a)\|_* - 1)^2]. \quad (17)$$

Remarks:

- In practice the distribution ν mimicks uniform (linearly interpolated) sampling as follows:

$$a \sim \text{Uniform}(a_i, h_x(\omega_i)).$$

- Spectral normalization: Divide each weight matrix by their spectral norm [20].

Practical implementation of GANs

Stochastic training of Wasserstein GANs

Input: primal and “dual” learning rates γ_t and α_m , primal iterations T , “dual” network d_y , generator network h_x , noise distribution p_Ω , real distribution $\hat{\mu}_n$, primal and dual batch sizes B, K , “dual” iterations M .

```
1. initialize  $\mathbf{x}^0$ 
2. For  $t = 0, 1, \dots, T - 1$ :
   For  $m = 0, 1, \dots, M - 1$ :
      initialize  $\mathbf{y}^0$ ,
      draw noise sample  $\omega_1, \dots, \omega_K \sim p_\Omega$ 
      draw real samples  $r_1, \dots, r_K \sim \hat{\mu}_n$ 
      “dual” pseudo-loss  $L(\mathbf{y}) := K^{-1} \sum_{i=1}^K d_y(r_i) - d_y(h_{\mathbf{x}^t}(\omega_i))$ 
      #update “dual” parameters  $\mathbf{y}^{m+1} = \mathbf{y}^m + \gamma_m \nabla_{\mathbf{y}} L(\mathbf{y}^m)$ 
      #enforce 1-Lipschitz constraint on  $d_y$ 
   end-For
   draw noise sample  $\omega_1, \dots, \omega_B \sim p_\Omega$ 
   generator pseudo-loss  $L(\mathbf{x}) := -B^{-1} \sum_{i=1}^B d_{\mathbf{y}^M}(h_{\mathbf{x}}(\omega_i))$ 
   update generator parameters  $\mathbf{x}^{t+1} = \mathbf{x}^t - \alpha_t \nabla_{\mathbf{x}} L(\mathbf{x}^t)$ 
end-For
```

#: Ideally, should be performed jointly.

Some historical background for a Turing award

Vanilla GAN [6]

$$\min_{\mathbf{x} \in \mathcal{X}} \max_{\mathbf{y} \in \mathcal{Y}} E_{\mathbf{a} \sim \hat{\mu}_n} [\log d_{\mathbf{y}}(\mathbf{a})] + E_{\omega \sim p_{\Omega}} [\log (1 - d_{\mathbf{y}}(h_{\mathbf{x}}(\omega)))] \quad (18)$$

- ▶ Binary cross-entropy modeling.
- ▶ $d_{\mathbf{y}}(\mathbf{a}) : \mathcal{Y} \rightarrow [0, 1]$ represents the probability that \mathbf{a} came from the real data distribution μ^{\natural} .

Observation: ◦ Minimizes Jensen-Shannon divergence:

$$\text{JSD}(\hat{\mu}_n \| h_{\mathbf{x}} \# p_{\Omega}) = \frac{1}{2} D(\hat{\mu}_n \| h_{\mathbf{x}} \# p_{\Omega}) + \frac{1}{2} D(h_{\mathbf{x}} \# p_{\Omega} \| \hat{\mu}_n).$$

Difficulties of GAN training






Figure: Mode collapse (left). Simultaneous vs alternating generator/discriminator updates (right).

- Heuristics galore!
- Difficult to enforce 1-Lipschitz constraint
- Overall a difficult minimax problem: Scalability, mode collapse, periodic cycling...
- Privacy concerns due to memorization


Application to 25 Gaussians: Algorithms matter [9]




(a) SGD



(b) Adam



(c) Mirror-GAN



(d) Mirror-Prox-GAN

Abstract minmax formulation

Minimax formulation

$$\min_{\mathbf{x} \in \mathcal{X}} \max_{\mathbf{y} \in \mathcal{Y}} \Phi(\mathbf{x}, \mathbf{y}), \quad (19)$$

where

- ▶ Φ is differentiable and nonconvex in \mathbf{x} and nonconcave in \mathbf{y} ,
- ▶ The domain is unconstrained, specifically $\mathcal{X} = \mathbb{R}^m$ and $\mathcal{Y} = \mathbb{R}^n$.

○ Key questions:

1. Where do the algorithms converge?
2. When do the algorithm converge?

Abstract minmax formulation

Minimax formulation

$$\min_{\mathbf{x} \in \mathcal{X}} \max_{\mathbf{y} \in \mathcal{Y}} \Phi(\mathbf{x}, \mathbf{y}), \quad (19)$$

where

- ▶ Φ is differentiable and nonconvex in \mathbf{x} and nonconcave in \mathbf{y} ,
- ▶ The domain is unconstrained, specifically $\mathcal{X} = \mathbb{R}^m$ and $\mathcal{Y} = \mathbb{R}^n$.

- Key questions:

1. Where do the algorithms converge?
2. When do the algorithm converge?

A buffet of negative results [2]

"Even when the objective is a Lipschitz and smooth differentiable function, deciding whether a min-max point exists, in fact even deciding whether an approximate min-max point exists, is NP-hard. More importantly, an approximate local min-max point of large enough approximation is guaranteed to exist, but finding one such point is PPAD-complete. The same is true of computing an approximate fixed point of the (Projected) Gradient Descent/Ascent update dynamics."

Solution concept

- Like for nonconvex problems in minimization we try to find a *local* solution.

Definition (Local Nash Equilibrium)

A pure strategy $(\mathbf{x}^*, \mathbf{y}^*)$ is called a Local Nash Equilibrium (LNE) if,

$$\Phi(\mathbf{x}^*, \mathbf{y}) \leq \Phi(\mathbf{x}^*, \mathbf{y}^*) \leq \Phi(\mathbf{x}, \mathbf{y}^*) \quad (\text{LNE})$$

for all \mathbf{x} and \mathbf{y} within some neighborhood of \mathbf{x}^* and \mathbf{y}^* , i.e., $\|\mathbf{x} - \mathbf{x}^*\| \leq \delta$ and $\|\mathbf{y} - \mathbf{y}^*\| \leq \delta$ for some $\delta > 0$.

Necessary conditions

Through a Taylor expansion around \mathbf{x}^* and \mathbf{y}^* one can show that a LNE implies,

$$\nabla_{\mathbf{x}} \Phi(\mathbf{x}, \mathbf{y}), -\nabla_{\mathbf{y}} \Phi(\mathbf{x}, \mathbf{y}) = 0$$

$$\nabla_{\mathbf{x}\mathbf{x}} \Phi(\mathbf{x}, \mathbf{y}), -\nabla_{\mathbf{y}\mathbf{y}} \Phi(\mathbf{x}, \mathbf{y}) \succeq 0$$




Figure: $\Phi(x, y) = x^2 - y^2$

Recall SGD results from Lecture 9

$$\min_{\mathbf{x}: \mathbf{x} \in \mathcal{X}} f(\mathbf{x})$$

- For a non-convex, smooth f , we have that
 1. SGD converges to the critical points of f as $N \rightarrow \infty$.
 2. SGD avoids strict saddles/traps ($\lambda_{\min}(\nabla^2 f(\mathbf{x}^*)) < 0$) almost surely.
 3. SGD remains close to Hurwicz minimizers (i.e., $\mathbf{x}^* : \lambda_{\min}(\nabla^2 f(\mathbf{x}^*)) > 0$ almost surely).

Recall SGD results from Lecture 9

$$\min_{\mathbf{x}: \mathbf{x} \in \mathcal{X}} f(\mathbf{x})$$

- For a non-convex, smooth f , we have that
 1. SGD converges to the critical points of f as $N \rightarrow \infty$.
 2. SGD avoids strict saddles/traps ($\lambda_{\min}(\nabla^2 f(\mathbf{x}^*)) < 0$) almost surely.
 3. SGD remains close to Hurwicz minimizers (i.e., $\mathbf{x}^* : \lambda_{\min}(\nabla^2 f(\mathbf{x}^*)) > 0$ almost surely).
- Nail in the coffin:
 - ▶ not even sure if we obtain stochastic descent directions by approximately solving inner problems in GANs.
 - ▶ GANs are fundamentally different from adversarial training!
- Need more direct approaches with the stochastic gradient estimates.

Basic algorithms for minimax

- Given $\min_{\mathbf{x} \in \mathcal{X}} \max_{\mathbf{y} \in \mathcal{Y}} \Phi(\mathbf{x}, \mathbf{y})$, define $V(\mathbf{z}) = [\nabla_{\mathbf{x}} \Phi(\mathbf{x}, \mathbf{y}), -\nabla_{\mathbf{y}} \Phi(\mathbf{x}, \mathbf{y})]$ with $\mathbf{z} = [\mathbf{x}, \mathbf{y}]$.




Figure: Trajectory of different algorithms for a simple bilinear game $\min_x \max_y xy$.

- (In)Famous algorithms
 - Gradient Descent Ascent (GDA)
 - Proximal point method (PPM)
 - Extra-gradient (EG)
 - Optimistic Gradient Descent Ascent (OGDA)
 - Reflected-Forward-Backward-Splitting (RFBS)
- EG and OGDA are approximations of the PPM
 - $\mathbf{z}^{k+1} = \mathbf{z}^k - \alpha V(\mathbf{z}^k)$.
 - $\mathbf{z}^{k+1} = \mathbf{z}^k - \alpha V(\mathbf{z}^{k+1})$.
 - $\mathbf{z}^{k+1} = \mathbf{z}^k - \alpha V(\mathbf{z}^k - \alpha V(\mathbf{z}^{k-1}))$
 - $\mathbf{z}^{k+1} = \mathbf{z}^k - \alpha [2V(\mathbf{z}^k) - V(\mathbf{z}^{k-1})]$
 - $\mathbf{z}^{k+1} = \mathbf{z}^k - \alpha V(2\mathbf{z}^k - \mathbf{z}^{k-1})$

Generalized Robbins-Monro schemes

- Given $\min_{\mathbf{x} \in \mathcal{X}} \max_{\mathbf{y} \in \mathcal{Y}} \Phi(\mathbf{x}, \mathbf{y})$, define $V(\mathbf{z}) = [\nabla_{\mathbf{x}} \Phi(\mathbf{x}, \mathbf{y}), -\nabla_{\mathbf{y}} \Phi(\mathbf{x}, \mathbf{y})]$ with $\mathbf{z} = [\mathbf{x}, \mathbf{y}]$.
- Given $V(\mathbf{z})$, define stochastic estimates of $V(\mathbf{z}, \zeta) = V(\mathbf{z}) + U(\mathbf{z}, \zeta)$, where
 - ▶ $U(\mathbf{z}, \zeta)$ is a bias term
 - ▶ We often have unbiasedness: $EU(\mathbf{z}, \zeta) = 0$
 - ▶ The bias term can have bounded moments
 - ▶ We often have bounded variance: $P(\|U(\mathbf{z}, \zeta)\| \geq t) \leq 2 \exp -\frac{t^2}{2\sigma^2}$ for $\sigma > 0$.
- An abstract template for generalized Robbins-Monro schemes, dubbed as \mathcal{A} :

$$\mathbf{z}^{k+1} = \mathbf{z}^k - \alpha_k V(\mathbf{z}^k, \zeta^k)$$

The dessert section in the buffet of negative results: [10]

1. Bounded trajectories of \mathcal{A} always converge to an internally chain-transitive (ICT) set.
2. Trajectories of \mathcal{A} may converge with arbitrarily high probability to spurious attractors that contain no critical point of Φ .

A deterministic, simple example beyond convex-concave

- Extragradient method: $\mathbf{z}^{k+1/2} = \mathbf{z}^k - \alpha_k V(\mathbf{z}^k)$, $\mathbf{z}^{k+1} = \mathbf{z}^k - \alpha_k V(\mathbf{z}^{k+1/2})$.

Example (Almost bilinear)

$$\Phi(\mathbf{x}, \mathbf{y}) = \mathbf{x}\mathbf{y} + \varepsilon\phi(\mathbf{y}) \quad (20)$$

where $\varepsilon > 0$ and $\phi(\mathbf{y}) = \frac{1}{2}\mathbf{y}^2 - \frac{1}{4}\mathbf{y}^4$.




Figure: Extra-gradient on (Almost bilinear) with $\varepsilon = 0.1$ converges to a stable limit cycle near an unstable critical point.

Example (Forsaken)

$$\Phi(\mathbf{x}, \mathbf{y}) = \mathbf{x}(\mathbf{y} - 0.1) + \phi(\mathbf{x}) - \phi(\mathbf{y}) \quad (21)$$

where $\phi(\mathbf{z}) = \frac{1}{4}\mathbf{z}^2 - \frac{1}{2}\mathbf{z}^4 + \frac{1}{6}\mathbf{z}^6$.




Figure: Extra-gradient on (Forsaken) can converge to a stable limit cycle. the white contour indicates the unstable limit cycle.

ExtraAdam

ExtraAdam for GANs [5]




Input. Step size γ , exponential decay rates $\eta_1, \eta_2 \in [0, 1)$

1. Set $\mathbf{m}_0, \mathbf{v}_0 = 0$
2. For $k = 0, 1, \dots$, iterate

$$\left\{ \begin{array}{lcl} \mathbf{g}_k & = V(\mathbf{z}^k, \zeta^k) \\ \mathbf{m}_{k-1/2} & = \eta_1 \mathbf{m}_{k-1} + (1 - \eta_1) \mathbf{g}_k & \leftarrow \text{1st order estimate} \\ \mathbf{v}_{k-1/2} & = \eta_2 \mathbf{v}_{k-1} + (1 - \eta_2) \mathbf{g}_k^2 & \leftarrow \text{2nd order estimate} \\ \hat{\mathbf{m}}_{k-1/2} & = \mathbf{m}_{k-1/2} / (1 - \eta_1^k) & \leftarrow \text{Bias correction} \\ \hat{\mathbf{v}}_{k-1/2} & = \mathbf{v}_{k-1/2} / (1 - \eta_2^k) & \leftarrow \text{Bias correction} \\ \mathbf{z}^{k+1/2} & = \mathbf{z}^k - \gamma \hat{\mathbf{m}}_{k-1/2} / (\sqrt{\hat{\mathbf{v}}_{k-1/2}} + \epsilon) & \leftarrow \text{Extrapolation step} \\ \mathbf{g}_{k+1/2} & = V(\mathbf{z}^{k+1/2}, \zeta^{k+1/2}) \\ \mathbf{m}_k & = \eta_1 \mathbf{m}_{k-1/2} + (1 - \eta_1) \mathbf{g}_{k+1/2} & \leftarrow \text{1st order estimate} \\ \mathbf{v}_k & = \eta_2 \mathbf{v}_{k-1/2} + (1 - \eta_2) \mathbf{g}_{k+1/2}^2 & \leftarrow \text{2nd order estimate} \\ \hat{\mathbf{m}}_k & = \mathbf{m}_k / (1 - \eta_1^k) & \leftarrow \text{Bias correction} \\ \hat{\mathbf{v}}_k & = \mathbf{v}_k / (1 - \eta_2^k) & \leftarrow \text{Bias correction} \\ \mathbf{z}^{k+1} & = \mathbf{z}^k - \gamma \hat{\mathbf{m}}_k / (\sqrt{\hat{\mathbf{v}}_k} + \epsilon) & \leftarrow \text{Update step} \end{array} \right.$$




Output : \mathbf{z}^k

Real LSUN Dataset: RMSProp, $4 \times 10^4, 8 \times 10^4, \times 10^5$ iterations [9]





(a) RMSProp

Real LSUN Dataset: Adam, $4 \times 10^4, 8 \times 10^4, \times 10^5$ iterations [9]



(b) Adam

Real LSUN Dataset: Mirror-GAN, $4 \times 10^4, 8 \times 10^4, \times 10^5$ iterations [9]



(c) Mirror-GAN, Algorithm 3

Real LSUN Dataset: Extra-Adam, $4 \times 10^4, 8 \times 10^4, \times 10^5$ iterations [9]



(d) Simultaneous Extra-Adam



(e) Alternated Extra-Adam

Wrap up!

1-Wasserstein Distributionally Robust Optimization (WDRO) [14]

Let W_1 be the 1-Wasserstein distance of probability measures over $\mathbb{R}^p \times \{\pm 1\}$ corresponding to the distance:

$$d((\mathbf{a}_1, \mathbf{b}_1), (\mathbf{a}_2, \mathbf{b}_2)) = \begin{cases} \|\mathbf{a}_1 - \mathbf{a}_2\|, & \text{if } \mathbf{b}_1 = \mathbf{b}_2; \\ \infty, & \text{otherwise.} \end{cases} \quad (22)$$

Let μ be a fixed probability distribution. The 1-Wasserstein Distributionally Robust Optimization Problem with radius ϵ is given by the following minimax formulation:

$$\text{WDRO}(\epsilon) = \min_{\mathbf{x}} \max_{\nu} \left\{ \mathbb{E}_{\mathbf{a}, \mathbf{b} \sim \nu} L(h_{\mathbf{x}}(\mathbf{a}), \mathbf{b}) : W_1(\nu, \mu) \leq \epsilon \right\}. \quad (23)$$

Lemma (Connection between WDRO and AT [24])


Let μ_n be the empirical measure supported on the input-label dataset $\{(\mathbf{a}_i, \mathbf{b}_i) : i = 1, \dots, n\}$. The $\text{WDRO}(\epsilon)$ problem is a relaxation (upper bound) of the adversarial training objective:

$$\min_{\mathbf{x}} \frac{1}{n} \sum_{i=1}^n \max_{\boldsymbol{\eta}: \|\boldsymbol{\eta}\| \leq \epsilon} L(h_{\mathbf{x}}(\mathbf{a}_i + \boldsymbol{\eta}), \mathbf{b}_i) \leq \text{WDRO}(\epsilon). \quad (24)$$

Wrap up!

- Continuing on Homework 2!

*Reinforcement Learning Game



- Environment: Markov Decision Process (MDP) $\mathcal{M} = (\mathcal{S}, \mathcal{A}, T, \gamma, P_0, R)$
- Agent: Parameterized deterministic policy $\mu_\theta : \mathcal{S} \rightarrow \mathcal{A}$, where $\theta \in \Theta$

Beyond supervised learning: Reinforcement Learning

At time step $t = 0$: $S_0 \sim P_0(\cdot)$

for $t = 1, 2, \dots$ do:

agent observes the environment's state $S_t \in \mathcal{S}$


agent chooses an action $A_t = \mu_\theta(S_t) \in \mathcal{A}$

agent receives a reward $R_{t+1} = R(S_t, A_t)$

agent finds itself in a new state $S_{t+1} \sim T(\cdot | S_t, A_t)$


*Exploration vs. Exploitation in RL

- Challenge: Exploration vs. exploitation!



*Exploration vs. Exploitation in RL

- Challenge: Exploration vs. exploitation!




- Objective (non-concave):

$$\max_{\theta \in \Theta} J(\theta) := \mathbb{E} \left[\sum_{t=1}^{\infty} \gamma^{t-1} R_t \mid \mu_{\theta}, \mathcal{M} \right]$$

- ▷ The environment only reveals the rewards after actions
- ▷ Exploitation: Maximize objective by choosing the appropriate action

*Exploration vs. Exploitation in RL

- Challenge: Exploration vs. exploitation!



- Objective (non-concave):

$$\max_{\theta \in \Theta} J(\theta) := \mathbb{E} \left[\sum_{t=1}^{\infty} \gamma^{t-1} R_t \mid \mu_{\theta}, \mathcal{M} \right]$$

- ▷ The environment only reveals the rewards after actions
- ▷ Exploitation: Maximize objective by choosing the appropriate action
- ▷ Exploration: Gather information on other actions

*Standard Reinforcement Learning

- Markov Decision Process (MDP): $\mathcal{M} = (\mathcal{S}, \mathcal{A}, T, \gamma, P_0, R)$
 - ▷ \mathcal{S} : state space
 - ▷ \mathcal{A} : action space
 - ▷ $T : \mathcal{S} \times \mathcal{S} \times \mathcal{A} \rightarrow [0, 1]$: state transition dynamics
 - ▷ $\gamma \in (0, 1)$: discounting factor
 - ▷ $P_0 : \mathcal{S} \rightarrow [0, 1]$: initial state distribution
 - ▷ $R : \mathcal{S} \times \mathcal{A} \rightarrow \mathbb{R}$: reward function
- Agent's (deterministic) policy: $\mu : \mathcal{S} \rightarrow \mathcal{A}$

Reinforcement Learning Game

for $t = 1, 2, \dots$ do:

agent observes the environment's state $S_t \in \mathcal{S}$

agent chooses an action $A_t = \mu(S_t) \in \mathcal{A}$

agent receives a reward $R_{t+1} = R(S_t, A_t)$, and finds itself in a new state S_{t+1}

*Standard Reinforcement Learning

- Discounted return:

$$Z = \sum_{t=1}^{\infty} \gamma^{t-1} R_t$$

- State and state-action value functions:

$$\begin{aligned} V^\mu(s) &:= \mathbb{E}[Z \mid S_1 = s; \mu, \mathcal{M}] \\ Q^\mu(s, a) &:= \mathbb{E}[Z \mid S_1 = s, A_1 = a; \mu, \mathcal{M}] \end{aligned}$$

- Performance objective:

$$\max_{\mu} J(\mu) := \mathbb{E}_{s \sim \mathcal{D}} [V^\mu(s)] = \mathbb{E}_{s \sim \mathcal{D}} [Q^\mu(s, \mu(s))]$$

*Deterministic Policy Gradient

- Deterministic policy parametrization:

$$\{\mu_\theta : \theta \in \Theta\}$$

- The off-policy performance objective:

$$\max_{\theta \in \Theta} J(\theta) := J(\mu_\theta) = \mathbb{E}_{s \sim \mathcal{D}} [Q^{\mu_\theta}(s, \mu_\theta(s))]$$

- The off-policy gradient:

[23]

$$\begin{aligned}\nabla_\theta J(\theta) &\approx \mathbb{E}_{s \sim \mathcal{D}} [\nabla_\theta \mu_\theta(s) \nabla_a Q^{\mu_\theta}(s, a)|_{a=\mu_\theta(s)}] \\ &\approx \frac{1}{N} \sum \nabla_a Q^\phi(s, a) \nabla_\theta \mu_\theta(s)\end{aligned}$$

- ▷ biased gradient estimate
- ▷ function approximation Q^ϕ for critic

*An optimization interpretation

- Objective (non-concave):

$$\max_{\theta \in \Theta} J(\theta) := \mathbb{E} \left[\sum_{t=1}^{\infty} \gamma^{t-1} R_t \mid \mu_{\theta}, \mathcal{M} \right]$$

- Exploitation: Progress in the gradient direction

$$\theta_{t+1} \leftarrow \theta_t + \eta_t \nabla_{\theta} \widehat{J}(\theta_t)$$

- Exploration: Add stochasticity while collecting the episodes

▷ noise injection in the action space

[23, 18]

$$a = \mu_{\theta}(s) + \mathcal{N}(0, \sigma^2 I)$$

▷ noise injection in the parameter space

[21]

$$\tilde{\theta} = \theta + \mathcal{N}(0, \sigma^2 I)$$

*Robust Reinforcement Learning

- Discounted return:

$$Z = \sum_{t=1}^{\infty} \gamma^{t-1} R_t$$

- State and state-action value functions:

$$\begin{aligned} V^\mu(s) &:= \mathbb{E}[Z \mid S_1 = s; \mu, \mathcal{M}] \\ Q^\mu(s, a) &:= \mathbb{E}[Z \mid S_1 = s, A_1 = a; \mu, \mathcal{M}] \end{aligned}$$

- Recall the standard performance objective: $J(\mu) := \mathbb{E}_{s \sim \mathcal{D}} [V^\mu(s)] = \mathbb{E}_{s \sim \mathcal{D}} [Q^\mu(s, \mu(s))]$

- An action robust formulation:

$$\max_{\mu} \mathbb{E}_{s \sim \mathcal{D}} \left[\max_{\nu \in \mathcal{N}} Q^\mu(s, \mu(s) + \nu) \right]$$

- See [11] for further details and results.

References I

- [1] Martin Arjovsky, Soumith Chintala, and Léon Bottou.
Wasserstein gan.
arXiv preprint arXiv:1701.07875, 2017.
- [2] Constantinos Daskalakis, Stratis Skoulakis, and Manolis Zampetakis.
The complexity of constrained min-max optimization.
arXiv preprint arXiv:2009.09623, 2020.
- [3] Richard Mansfield Dudley.
The speed of mean glivenko-cantelli convergence.
The Annals of Mathematical Statistics, 40(1):40–50, 1969.
- [4] Marwa El Halabi.
Learning with structured sparsity: From discrete to convex and back.
Technical report, EPFL, 2018.
- [5] Gauthier Gidel, Hugo Berard, Gaëtan Vignoud, Pascal Vincent, and Simon Lacoste-Julien.
A variational inequality perspective on generative adversarial networks.
In *International Conference on Learning Representations*, 2018.

References II

- [6] Ian Goodfellow, Jean Pouget-Abadie, Mehdi Mirza, Bing Xu, David Warde-Farley, Sherjil Ozair, Aaron Courville, and Yoshua Bengio.
Generative adversarial nets.
In Z. Ghahramani, M. Welling, C. Cortes, N. D. Lawrence, and K. Q. Weinberger, editors, *Advances in Neural Information Processing Systems 27*, pages 2672–2680. Curran Associates, Inc., 2014.
- [7] Ian J Goodfellow, Jonathon Shlens, and Christian Szegedy.
Explaining and harnessing adversarial examples.
arXiv preprint arXiv:1412.6572, 2014.
- [8] Ishaan Gulrajani, Faruk Ahmed, Martin Arjovsky, Vincent Dumoulin, and Aaron C Courville.
Improved training of wasserstein gans.
In *Advances in neural information processing systems*, pages 5767–5777, 2017.
- [9] Ya-Ping Hsieh, Chen Liu, and Volkan Cevher.
Finding mixed Nash equilibria of generative adversarial networks.
volume 97 of *Proceedings of Machine Learning Research*, pages 2810–2819, Long Beach, California, USA, 09–15 Jun 2019. PMLR.
- [10] Ya-Ping Hsieh, Panayotis Mertikopoulos, and Volkan Cevher.
The limits of min-max optimization algorithms: convergence to spurious non-critical sets.
arXiv preprint arXiv:2006.09065, 2020.

References III

- [11] Parameswaran Kamalaruban, Yu-Ting Huang, Ya-Ping Hsieh, Paul Rolland, Cheng Shi, and Volkan Cevher.
Robust reinforcement learning via adversarial training with langevin dynamics.
arXiv preprint arXiv:2002.06063, 2020.
- [12] Tero Karras, Timo Aila, Samuli Laine, and Jaakko Lehtinen.
Progressive growing of GANs for improved quality, stability, and variation.
In *International Conference on Learning Representations*, 2018.
- [13] Ziko Kolter and Aleksander Madry.
Adversarial robustness - theory and practice.
NeurIPS 2018 tutorial: <https://adversarial-ml-tutorial.org/>.
- [14] Daniel Kuhn, Peyman Mohajerin Esfahani, Viet Anh Nguyen, and Soroosh Shafieezadeh-Abadeh.
Wasserstein Distributionally Robust Optimization: Theory and Applications in Machine Learning.
arXiv e-prints, page arXiv:1908.08729, August 2019.
- [15] Alexey Kurakin, Ian Goodfellow, and Samy Bengio.
Adversarial examples in the physical world.
arXiv preprint arXiv:1607.02533, 2016.

References IV

- [16] Alexey Kurakin, Ian Goodfellow, and Samy Bengio.
Adversarial machine learning at scale.
arXiv preprint arXiv:1611.01236, 2016.
- [17] Fabian Latorre, Paul Rolland, and Volkan Cevher.
Lipschitz constant estimation of neural networks via sparse polynomial optimization.
In *International Conference on Learning Representations*, 2020.
- [18] Timothy P Lillicrap, Jonathan J Hunt, Alexander Pritzel, Nicolas Heess, Tom Erez, Yuval Tassa, David Silver, and Daan Wierstra.
Continuous control with deep reinforcement learning.
arXiv preprint arXiv:1509.02971, 2015.
- [19] Aleksander Madry, Aleksandar Makelov, Ludwig Schmidt, Dimitris Tsipras, and Adrian Vladu.
Towards deep learning models resistant to adversarial attacks.
In *International Conference on Learning Representations*, 2018.
- [20] Takeru Miyato, Toshiki Kataoka, Masanori Koyama, and Yuichi Yoshida.
Spectral normalization for generative adversarial networks.
arXiv preprint arXiv:1802.05957, 2018.

References V

- [21] Matthias Plappert, Rein Houthooft, Prafulla Dhariwal, Szymon Sidor, Richard Y Chen, Xi Chen, Tamim Asfour, Pieter Abbeel, and Marcin Andrychowicz.
Parameter space noise for exploration.
arXiv preprint arXiv:1706.01905, 2017.
- [22] Kevin Scaman and Aladin Virmaux.
Lipschitz regularity of deep neural networks: analysis and efficient estimation.
arXiv e-prints, page arXiv:1805.10965, May 2018.
- [23] David Silver, Guy Lever, Nicolas Heess, Thomas Degris, Daan Wierstra, and Martin Riedmiller.
Deterministic policy gradient algorithms.
In *ICML*, 2014.
- [24] Matthew Staib and Stefanie Jegelka.
Distributionally robust deep learning as a generalization of adversarial training.
- [25] Mukund Sundararajan, Ankur Taly, and Qiqi Yan.
Axiomatic attribution for deep networks.
arXiv preprint arXiv:1703.01365, 2017.

References VI

- [26] Christian Szegedy, Wojciech Zaremba, Ilya Sutskever, Joan Bruna, Dumitru Erhan, Ian Goodfellow, and Rob Fergus.
Intriguing properties of neural networks.
arXiv preprint arXiv:1312.6199, 2013.
- [27] Jonathan Weed, Francis Bach, et al.
Sharp asymptotic and finite-sample rates of convergence of empirical measures in wasserstein distance.
Bernoulli, 25(4A):2620–2648, 2019.