

Mathematics of Data: From Theory to Computation

Prof. Volkan Cevher
volkan.cevher@epfl.ch

Lecture 13: Primal-dual optimization III: Lagrangian gradient methods

Laboratory for Information and Inference Systems (LIONS)
École Polytechnique Fédérale de Lausanne (EPFL)

EE-556 (Fall 2021)

lions@epfl



Google AI



EPFL

License Information for Mathematics of Data Slides

- ▶ This work is released under a [Creative Commons License](#) with the following terms:
- ▶ **Attribution**
 - ▶ The licensor permits others to copy, distribute, display, and perform the work. In return, licensees must give the original authors credit.
- ▶ **Non-Commercial**
 - ▶ The licensor permits others to copy, distribute, display, and perform the work. In return, licensees may not use the work for commercial purposes – unless they get the licensor's permission.
- ▶ **Share Alike**
 - ▶ The licensor permits others to distribute derivative works only under a license identical to the one that governs the licensor's work.
- ▶ [Full Text of the License](#)

Recall: Swiss army knife of convex formulations

A primal problem prototype

$$f^* := \min_{\mathbf{x} \in \mathbb{R}^p} \left\{ f(\mathbf{x}) : \mathbf{A}\mathbf{x} - \mathbf{b} \in \mathcal{K}, \mathbf{x} \in \mathcal{X} \right\},$$

- ▶ f is proper, closed and **convex**
- ▶ \mathcal{X} and \mathcal{K} are nonempty, closed **convex** sets
- ▶ $\mathbf{A} \in \mathbb{R}^{n \times p}$ and $\mathbf{b} \in \mathbb{R}^n$ are known
- ▶ An optimal solution \mathbf{x}^* satisfies $f(\mathbf{x}^*) = f^*$, $\mathbf{A}\mathbf{x}^* - \mathbf{b} \in \mathcal{K}$ and $\mathbf{x}^* \in \mathcal{X}$

Broad context for the problem template:

- ▶ Many **real-world applications** (e.g., linear inverse problems) can be directly formulated as (3).
- ▶ Often times, computational limitations require the translation of existing unconstrained problems (e.g., composite convex minimization, consensus optimization, and convex splitting) into constrained ones (3).
- ▶ Many **standard convex optimization** formulations naturally fall under (3), such as *linear programming*, *convex quadratic programming*, *second order cone programming*, *semidefinite programming* and *geometric programming*.

Recall - Swiss army knife of convex formulations

A primal problem prototype

$$f^* := \min_{\mathbf{x} \in \mathbb{R}^p} \left\{ f(\mathbf{x}) : \mathbf{A}\mathbf{x} - \mathbf{b} \in \mathcal{K}, \mathbf{x} \in \mathcal{X} \right\},$$

- ▶ f is proper, closed and convex
- ▶ \mathcal{X} and \mathcal{K} are nonempty, closed convex sets
- ▶ $\mathbf{A} \in \mathbb{R}^{n \times p}$ and $\mathbf{b} \in \mathbb{R}^n$ are known
- ▶ An optimal solution \mathbf{x}^* to (3) satisfies $f(\mathbf{x}^*) = f^*$, $\mathbf{A}\mathbf{x}^* - \mathbf{b} \in \mathcal{K}$ and $\mathbf{x}^* \in \mathcal{X}$

A simplified template

$$f^* := \min_{\mathbf{x} \in \mathbb{R}^p} \left\{ f(\mathbf{x}) : \mathbf{A}\mathbf{x} = \mathbf{b}, \right\}, \quad (1)$$

- ▶ f is proper, closed and convex
- ▶ $\mathbf{A} \in \mathbb{R}^{n \times p}$ and $\mathbf{b} \in \mathbb{R}^n$ are known
- ▶ An optimal solution \mathbf{x}^* to (1) satisfies $f(\mathbf{x}^*) = f^*$, $\mathbf{A}\mathbf{x}^* = \mathbf{b}$.

Recall: Finding the solutions in affine constrained convex minimization

A performance metric: Time-to-reach ϵ

time-to-reach ϵ = number of iterations to reach ϵ \times per iteration time

A key issue: Number of iterations to reach ϵ

The notion of ϵ -accuracy is elusive in constrained optimization!

Our definition of ϵ -accurate solutions [36]

Given a numerical tolerance $\epsilon \geq 0$, a point $\mathbf{x}_\epsilon^* \in \mathbb{R}^p$ is called an ϵ -solution of (1) if

$$\begin{cases} f(\mathbf{x}_\epsilon^*) - f^* \leq \epsilon & (\text{objective residual}), \\ \|\mathbf{Ax}_\epsilon^* - \mathbf{b}\| \leq \epsilon & (\text{feasibility gap}), \end{cases}$$

- When \mathbf{x}^* is unique, we can also obtain $\|\mathbf{x}_\epsilon^* - \mathbf{x}^*\| \leq \epsilon$ (iterate residual).

Remark:

- ϵ can be different for the objective, feasibility gap, or the iterate residual.

Plenty of primal-dual methods for solving (1):

- **Penalty and augmented Lagrangian methods:**

- ▶ Exact penalty method [3].
- ▶ Quadratic penalty method [4].
- ▶ Augmented Lagrangian method [23, 30].

See Lecture 12
This lecture

- **Variants of the Arrow-Hurwitz's method:**

- ▶ Proximal-based decomposition (Chen-Teboulle's algorithm) [9].
- ▶ Primal-dual Hybrid Gradient (PDHG) method and its variants [15, 18].
- ▶ Chambolle-Pock's algorithm [7], and its variants, e.g., He-Yuan's variant [20].

See Lecture 12

- **Splitting techniques from monotone inclusions:**

- ▶ Primal-dual splitting algorithms [2, 10, 37, 11, 12].
- ▶ Three-operator splitting [13].

See Lecture 12

- **Dual splitting techniques:**

- ▶ Alternating minimization algorithms (AMA) [16, 37].
- ▶ Alternating direction methods of multipliers (ADMM) [14, 22].
- ▶ Accelerated variants of AMA and ADMM [12, 19].
- ▶ Preconditioned ADMM, Linearized ADMM and inexact Uzawa algorithms [7, 27].

- **Second-order decomposition methods:**

- ▶ Dual (quasi) Newton methods [38].
- ▶ Smoothing decomposition methods via barriers functions [26, 34].

Recall: Quadratic penalty & Lagrangian formulations

- **The problem:** $f^* := \min_{\mathbf{x} \in \mathbb{R}^p} \left\{ f(\mathbf{x}) : \mathbf{A}\mathbf{x} = \mathbf{b} \right\}$

- **Reformulations:**

| Quadratic Penalty | The Lagrangian |
|--|--|
| $f^* = f(\mathbf{x}^*) + \frac{\beta}{2} \ \mathbf{A}\mathbf{x}^* - \mathbf{b}\ ^2, \quad \forall \beta > 0.$ | $f^* = f(\mathbf{x}^*) + \max_{\boldsymbol{\lambda} \in \mathbb{R}^n} \langle \boldsymbol{\lambda}, \mathbf{A}\mathbf{x}^* - \mathbf{b} \rangle.$ |
| $F_\beta(\mathbf{x}) = f(\mathbf{x}) + \frac{\beta}{2} \ \mathbf{A}\mathbf{x} - \mathbf{b}\ ^2.$ | $F_{\boldsymbol{\lambda}}(\mathbf{x}) = f(\mathbf{x}) + \max_{\boldsymbol{\lambda} \in \mathbb{R}^n} \langle \boldsymbol{\lambda}, \mathbf{A}\mathbf{x} - \mathbf{b} \rangle$ $= f(\mathbf{x}) + \begin{cases} 0, & \text{if } \mathbf{A}\mathbf{x} = \mathbf{b}, \\ +\infty, & \text{if } \mathbf{A}\mathbf{x} \neq \mathbf{b}. \end{cases}$ |
| $\min_{\mathbf{x} \in \mathbb{R}^p} \left\{ f(\mathbf{x}) : \mathbf{A}\mathbf{x} = \mathbf{b} \right\} \equiv \lim_{\beta \rightarrow \infty} \min_{\mathbf{x} \in \mathbb{R}^p} \left\{ f(\mathbf{x}) + \frac{\beta}{2} \ \mathbf{A}\mathbf{x} - \mathbf{b}\ ^2 \right\}$ | $\min_{\mathbf{x} \in \mathbb{R}^p} \left\{ f(\mathbf{x}) : \mathbf{A}\mathbf{x} = \mathbf{b} \right\} \equiv \min_{\mathbf{x} \in \mathbb{R}^p} \max_{\boldsymbol{\lambda} \in \mathbb{R}^n} \left\{ f(\mathbf{x}) + \langle \boldsymbol{\lambda}, \mathbf{A}\mathbf{x} - \mathbf{b} \rangle \right\}$ |

Recall: Quadratic penalty & Lagrangian methods

o **The problem:** $f^* := \min_{\mathbf{x} \in \mathbb{R}^p} \left\{ f(\mathbf{x}) : \mathbf{Ax} = \mathbf{b} \right\}$

o **The methods:**

| Quadratic penalty method (QP) | Dual subgradient method (DSGM) |
|--|---|
| <p>1. Choose $\mathbf{x}^0 \in \mathbb{R}^p$ and $\beta_0 > 0$.</p> <p>2. For $k = 0, 1, \dots$, perform:</p> <p>2.a. $\mathbf{x}^k := \arg \min_{\mathbf{x} \in \mathbb{R}^p} \left\{ f(\mathbf{x}) + \frac{\beta_k}{2} \ \mathbf{Ax} - \mathbf{b}\ ^2 \right\}$.</p> <p>2.b. Update $\beta_{k+1} > \beta_k$.</p> <p>o Drawbacks:</p> <ul style="list-style-type: none">► $\mathbf{x}^k := \arg \min_{\mathbf{x} \in \mathbb{R}^p} \left\{ f(\mathbf{x}) + \frac{\beta_k}{2} \ \mathbf{Ax} - \mathbf{b}\ ^2 \right\}$ becomes ill-conditioned as $\beta_k \rightarrow \infty$. | <p>1. Choose $\boldsymbol{\lambda}^0 \in \mathbb{R}^n$.</p> <p>2. For $k = 0, 1, \dots$, perform:</p> <p>2.a. $\mathbf{x}^*(\boldsymbol{\lambda}^k) := \arg \min_{\mathbf{x} \in \mathbb{R}^p} \left\{ \mathcal{L}(\mathbf{x}, \boldsymbol{\lambda}^k) := f(\mathbf{x}) + \langle \boldsymbol{\lambda}^k, \mathbf{Ax} - \mathbf{b} \rangle \right\}$.</p> <p>2.b. Compute the subgradient $\nabla d(\boldsymbol{\lambda}^k) := \mathbf{Ax}^*(\boldsymbol{\lambda}^k) - \mathbf{b}$.</p> <p>2.c. Update $\boldsymbol{\lambda}^{k+1} := \boldsymbol{\lambda}^k + \frac{R}{\sqrt{k+1}} \nabla d(\boldsymbol{\lambda}^k)$, where R is a given constant.</p> <p>o Drawbacks:</p> <ul style="list-style-type: none">► $d(\boldsymbol{\lambda})$ is not necessarily smooth \implies slower rates.► $\mathbf{x}^*(\boldsymbol{\lambda}^k)$ is not necessarily well-defined for all $\boldsymbol{\lambda}$.► Finding R is not always straightforward. |

Unifying the Lagrangian and the penalty approaches

- Quadratic penalty:

$$F_\beta(\mathbf{x}) = f(\mathbf{x}) + \frac{\beta}{2} \|\mathbf{Ax} - \mathbf{b}\|^2$$

+

- The Lagrangian:

$$\mathcal{L}(\mathbf{x}, \boldsymbol{\lambda}) = f(\mathbf{x}) + \langle \boldsymbol{\lambda}, \mathbf{Ax} - \mathbf{b} \rangle$$

↓

- Augmented Lagrangian (AL):** $\mathcal{L}_\beta(\mathbf{x}, \boldsymbol{\lambda}) = f(\mathbf{x}) + \langle \boldsymbol{\lambda}, \mathbf{Ax} - \mathbf{b} \rangle + \frac{\beta}{2} \|\mathbf{Ax} - \mathbf{b}\|^2$

Properties of AL

- The dual function is concave and $\frac{1}{\beta}$ -smooth:

$$d_\beta(\boldsymbol{\lambda}) = \min_{\mathbf{x} \in \mathbb{R}^p} \left\{ f(\mathbf{x}) + \langle \boldsymbol{\lambda}, \mathbf{Ax} - \mathbf{b} \rangle + \frac{\beta}{2} \|\mathbf{Ax} - \mathbf{b}\|^2 \right\}.$$

Can apply gradient or accelerated gradient methods in the dual!

- β does not need to increase until infinity.

No more ill-conditioned subproblems!


Example: Behavior of the AL dual function

Consider a constrained convex problem:

$$\begin{aligned} \min_{\mathbf{x} \in \mathbb{R}^3} \quad & \left\{ f(\mathbf{x}) := x_1^2 + x_2^2 \right\}, \\ \text{s.t.} \quad & 2x_3 - x_1 - x_2 = 1, \\ & \mathbf{x} \in \mathcal{X} := [-2, 2] \times [-2, 2] \times [0, 2]. \end{aligned}$$

The **AL dual function** is **concave, smooth** and defined as

$$d_\beta(\lambda) := \min_{\mathbf{x} \in \mathcal{X}} \left\{ x_1^2 + x_2^2 + \lambda(2x_3 - x_1 - x_2 - 1) + (\beta/2)\|2x_3 - x_1 - x_2 - 1\|_2^2 \right\}$$




Example: Behavior of the AL dual function

Consider a constrained convex problem:


$$\begin{aligned} \min_{\mathbf{x} \in \mathbb{R}^3} \quad & \left\{ f(\mathbf{x}) := x_1^2 + x_2^2 \right\}, \\ \text{s.t.} \quad & 2x_3 - x_1 - x_2 = 1, \\ & \mathbf{x} \in \mathcal{X} := [-2, 2] \times [-2, 2] \times [0, 2]. \end{aligned}$$

The **AL dual function** is **concave, smooth** and defined as

$$d_\beta(\lambda) := \min_{\mathbf{x} \in \mathcal{X}} \left\{ x_1^2 + x_2^2 + \lambda(2x_3 - x_1 - x_2 - 1) + (\beta/2)\|2x_3 - x_1 - x_2 - 1\|_2^2 \right\}$$



VS



Augmented dual problem

- Dual problem:

$$d^* := \max_{\boldsymbol{\lambda} \in \mathbb{R}^n} \left\{ d(\boldsymbol{\lambda}) = \min_{\mathbf{x} \in \mathbb{R}^p} f(\mathbf{x}) + \langle \boldsymbol{\lambda}, \mathbf{A}\mathbf{x} - \mathbf{b} \rangle \right\}. \quad (2)$$

- Augmented dual problem:

$$d_\beta^* := \max_{\boldsymbol{\lambda} \in \mathbb{R}^n} \left\{ d_\beta(\boldsymbol{\lambda}) = \min_{\mathbf{x} \in \mathbb{R}^p} f(\mathbf{x}) + \langle \boldsymbol{\lambda}, \mathbf{A}\mathbf{x} - \mathbf{b} \rangle + \frac{\beta}{2} \|\mathbf{A}\mathbf{x} - \mathbf{b}\|^2 \right\}, \quad \beta > 0. \quad (3)$$

Augmented dual problem

- Dual problem:

$$d^* := \max_{\lambda \in \mathbb{R}^n} \left\{ d(\lambda) = \min_{x \in \mathbb{R}^p} f(x) + \langle \lambda, Ax - b \rangle \right\}. \quad (2)$$

- Augmented dual problem:

$$d_\beta^* := \max_{\lambda \in \mathbb{R}^n} \left\{ d_\beta(\lambda) = \min_{x \in \mathbb{R}^p} f(x) + \langle \lambda, Ax - b \rangle + \frac{\beta}{2} \|Ax - b\|^2 \right\}, \quad \beta > 0. \quad (3)$$

Relation between augmented dual problem and dual problem

If a primal solution exists and [Slater's condition](#) holds, we have

- ▶ The [dual solution set](#) of (3) coincides with the [one](#) of the dual problem (2).
- ▶ $f^* = d^* = d_\beta^*$ for any $\beta > 0$.

- Recall: The augmented dual problem (3) is [smooth](#) and [concave](#)

⇒ **Gradient and accelerated gradient methods** can be applied to solve it.

Augmented Lagrangian method: The ideal algorithm

$$d_\beta(\boldsymbol{\lambda}) = \min_{\mathbf{x} \in \mathbb{R}^p} \left\{ f(\mathbf{x}) + \langle \boldsymbol{\lambda}, \mathbf{A}\mathbf{x} - \mathbf{b} \rangle + \frac{\beta}{2} \|\mathbf{A}\mathbf{x} - \mathbf{b}\|^2 \right\} \quad (4)$$
$$\mathbf{x}_\beta^*(\boldsymbol{\lambda}) \in \arg \min_{\mathbf{x} \in \mathbb{R}^p} \left\{ f(\mathbf{x}) + \langle \boldsymbol{\lambda}, \mathbf{A}\mathbf{x} - \mathbf{b} \rangle + \frac{\beta}{2} \|\mathbf{A}\mathbf{x} - \mathbf{b}\|^2 \right\}$$

Augmented Lagrangian method (ALM)

1. Choose $\boldsymbol{\lambda}^0 \in \mathbb{R}^n$ and $\beta > 0$.
 2. For $k = 0, 1, \dots$:
 - 2.a. Solve (4).
 - 2.b. Compute $\nabla d_\beta(\boldsymbol{\lambda}^k) := \mathbf{A}\mathbf{x}_\beta^*(\boldsymbol{\lambda}^k) - \mathbf{b}$.
 - 2.c. Update $\boldsymbol{\lambda}^{k+1} := \boldsymbol{\lambda}^k + \beta \nabla d_\beta(\boldsymbol{\lambda}^k)$.
-

Augmented Lagrangian method: The ideal algorithm

$$d_\beta(\boldsymbol{\lambda}) = \min_{\mathbf{x} \in \mathbb{R}^p} \left\{ f(\mathbf{x}) + \langle \boldsymbol{\lambda}, \mathbf{A}\mathbf{x} - \mathbf{b} \rangle + \frac{\beta}{2} \|\mathbf{A}\mathbf{x} - \mathbf{b}\|^2 \right\} \quad (4)$$
$$\mathbf{x}_\beta^*(\boldsymbol{\lambda}) \in \arg \min_{\mathbf{x} \in \mathbb{R}^p} \left\{ f(\mathbf{x}) + \langle \boldsymbol{\lambda}, \mathbf{A}\mathbf{x} - \mathbf{b} \rangle + \frac{\beta}{2} \|\mathbf{A}\mathbf{x} - \mathbf{b}\|^2 \right\}$$

| Augmented Lagrangian method (ALM) | Accelerated ALM (AALM) |
|--|---|
| <ol style="list-style-type: none">1. Choose $\boldsymbol{\lambda}^0 \in \mathbb{R}^n$ and $\beta > 0$.2. For $k = 0, 1, \dots$:<ol style="list-style-type: none">2.a. Solve (4).2.b. Compute $\nabla d_\beta(\boldsymbol{\lambda}^k) := \mathbf{A}\mathbf{x}_\beta^*(\boldsymbol{\lambda}^k) - \mathbf{b}$.2.c. Update $\boldsymbol{\lambda}^{k+1} := \boldsymbol{\lambda}^k + \beta \nabla d_\beta(\boldsymbol{\lambda}^k)$. | <ol style="list-style-type: none">1. Choose $\boldsymbol{\lambda}^0 \in \mathbb{R}^n$ and $\beta > 0$. Set $\tilde{\boldsymbol{\lambda}}^0 := \boldsymbol{\lambda}^0$ and $t_0 := 1$2. For $k = 0, 1, \dots$, perform:<ol style="list-style-type: none">2.a. Solve (4).2.b. Compute $\nabla d_\beta(\tilde{\boldsymbol{\lambda}}^k) := \mathbf{A}\mathbf{x}_\beta^*(\tilde{\boldsymbol{\lambda}}^k) - \mathbf{b}$.2.c. Update $\boldsymbol{\lambda}^{k+1} := \tilde{\boldsymbol{\lambda}}_k + \beta \nabla d_\beta(\tilde{\boldsymbol{\lambda}}^k)$, $\tilde{\boldsymbol{\lambda}}^{k+1} := \boldsymbol{\lambda}^{k+1} + ((t_k - 1)/t_{k+1})(\boldsymbol{\lambda}^{k+1} - \boldsymbol{\lambda}^k),$$t_{k+1} := (1 + \sqrt{1 + 4t_k^2})/2.$ |

Convergence of ALM and AALM

Theorem (Convergence [21])

- Let $\{\boldsymbol{\lambda}^k\}$ be the sequence generated by ALM. Then

$$d^* - d_\beta(\boldsymbol{\lambda}^k) \leq \frac{\|\boldsymbol{\lambda}^0 - \boldsymbol{\lambda}^*\|_2^2}{2\beta(k+1)}.$$

- Let $\{\boldsymbol{\lambda}^k\}$ be the sequence generated by AALM. Then

$$d^* - d_\beta(\boldsymbol{\lambda}^k) \leq \frac{2\|\boldsymbol{\lambda}^0 - \boldsymbol{\lambda}^*\|_2^2}{\beta(k+1)^2}.$$

Remarks:

- Guarantees are given for the dual problem and not for the primal!
- Approximate solution for primal via averaging: $\mathbf{x}^\epsilon = \frac{1}{k} \sum_{i=0}^{k-1} \mathbf{x}_\beta^*(\boldsymbol{\lambda}^i)$ [44]

Drawbacks and enhancements

- At each step, ALM solves

$$\mathbf{x}_\beta^*(\boldsymbol{\lambda}) := \arg \min_{\mathbf{x} \in \mathbb{R}^p} \left\{ \mathcal{L}_\beta(\mathbf{x}, \boldsymbol{\lambda}) := f(\mathbf{x}) + \langle \boldsymbol{\lambda}, \mathbf{A}\mathbf{x} - \mathbf{b} \rangle + \frac{\beta}{2} \|\mathbf{A}\mathbf{x} - \mathbf{b}\|^2 \right\}. \quad (5)$$

Drawbacks

1. **Drawback 1:** The quadratic term $\|\mathbf{A}\mathbf{x} - \mathbf{b}\|^2$ in (5) **destroys** the **separability** as well as the **tractable proximity** of f .
2. **Drawback 2:** Solving (5) exactly is **impractical**.

Drawbacks and enhancements

- At each step, ALM solves

$$\mathbf{x}_\beta^*(\boldsymbol{\lambda}) := \arg \min_{\mathbf{x} \in \mathbb{R}^p} \left\{ \mathcal{L}_\beta(\mathbf{x}, \boldsymbol{\lambda}) := f(\mathbf{x}) + \langle \boldsymbol{\lambda}, \mathbf{A}\mathbf{x} - \mathbf{b} \rangle + \frac{\beta}{2} \|\mathbf{A}\mathbf{x} - \mathbf{b}\|^2 \right\}. \quad (5)$$

Drawbacks

1. **Drawback 1:** The quadratic term $\|\mathbf{A}\mathbf{x} - \mathbf{b}\|^2$ in (5) **destroys** the **separability** as well as the **tractable proximity** of f .
2. **Drawback 2:** Solving (5) exactly is **impractical**.

Enhancements

1. Allow **inexactness** of solving (5), while guaranteeing the **same convergence rate**.
2. Linearize the term $\|\mathbf{A}\mathbf{x} - \mathbf{b}\|^2$ in the same way we did for Quadratic Penalty formulations.

An inexact approach for subproblems of ALM

- Primal subproblem as a **composite optimization problem**:

$$\mathbf{x}_\beta^*(\boldsymbol{\lambda}) := \arg \min_{\mathbf{x} \in \mathbb{R}^p} \left\{ \mathcal{L}_\beta(\mathbf{x}, \boldsymbol{\lambda}) := \underbrace{f(\mathbf{x}) + \langle \boldsymbol{\lambda}, \mathbf{A}\mathbf{x} - \mathbf{b} \rangle}_{=: h(x)} + \frac{\beta}{2} \underbrace{\|\mathbf{A}\mathbf{x} - \mathbf{b}\|^2}_{=: g(x)} \right\}. \quad (6)$$

proximally
tractable

⇒ can use **accelerated proximal methods** (e.g. FISTA) to solve this up to some accuracy.

Conceptual inexact augmented Lagrangian method:

- Choose $\boldsymbol{\lambda}^0 \in \mathbb{R}^n$, $\beta > 0$ and a decreasing sequence $\epsilon_k \geq 0$, $\forall k$.
- For $k = 0, 1, \dots$, perform:
 - Solve (6) with FISTA until $\mathcal{L}_\beta(\mathbf{x}_\beta^{\epsilon_k}(\boldsymbol{\lambda}^k), \boldsymbol{\lambda}^k) \leq \mathcal{L}_\beta(\mathbf{x}_\beta^*(\boldsymbol{\lambda}^k), \boldsymbol{\lambda}^k) + \epsilon_k$.
 - Update $\boldsymbol{\lambda}^{k+1} := \boldsymbol{\lambda}^k + \beta(\mathbf{A}\mathbf{x}_\beta^{\epsilon_k}(\boldsymbol{\lambda}^k) - \mathbf{b})$.

Remarks:

- Conceptual since $\mathbf{x}_\beta^*(\boldsymbol{\lambda}^k)$ is unknown.
- Solve (6) for increasing (**explicit**) number of iterations $m_k > 0$.
- See advanced material at the end of the lecture for DL-ASGARD method.

Linearized Augmented Lagrangian method (LALM)

1. Majorize the augmented Lagrangian:

$$\mathbf{x}^{k+1} := \arg \min_{\mathbf{x} \in \mathcal{X}} \left\{ f(\mathbf{x}) + \langle \boldsymbol{\lambda}, \mathbf{Ax} - \mathbf{b} \rangle + \frac{\beta}{2} \|\mathbf{Ax} - \mathbf{b}\|^2 + \frac{1}{2} \|\mathbf{x} - \mathbf{x}^k\|_{\mathbf{Q}_k}^2 \right\}.$$

2. Using the same calculation as in Lecture 12, when $\mathbf{Q}_k = \alpha_k \mathbf{I} - \beta \mathbf{A}^\top \mathbf{A} \succeq 0$ and $\alpha_k \geq \beta \|\mathbf{A}\|^2$, we get:

$$\mathbf{x}^{k+1} = \text{prox}_{\frac{1}{\alpha_k} f} \left(\mathbf{x}^k - \frac{1}{\alpha_k} \mathbf{A}^\top (\boldsymbol{\lambda}^k + \beta (\mathbf{Ax}^k - \mathbf{b})) \right)$$

3. Picking $\alpha_k = \beta \|\mathbf{A}\|^2$, we obtain the following method:

Accelerated LALM (Alg.1 + parameters of eq. (30) in [39])

1. Choose $\mathbf{x}^0 \in \mathbb{R}^p$, $\boldsymbol{\lambda}^0 \in \mathbb{R}^n$ and $\beta > 0$.

2. For $k = 0, 1, \dots$:

$$\mathbf{x}^{k+1} := \text{prox}_{\frac{1}{\beta \|\mathbf{A}\|^2} f} \left(\mathbf{x}^k - \frac{1}{\beta \|\mathbf{A}\|^2} \mathbf{A}^\top (\boldsymbol{\lambda}^k + \beta (\mathbf{Ax}^k - \mathbf{b})) \right),$$

$$\boldsymbol{\lambda}^{k+1} := \boldsymbol{\lambda}^k + \beta (\mathbf{Ax}^{k+1} - \mathbf{b}).$$

Convergence of Accelerated LALM

Theorem (Convergence result of Theorem 2.5 in [39])

Let $\beta > 0$ and define $\bar{\mathbf{x}}^k = \frac{1}{k} \sum_{i=1}^k \mathbf{x}^i$. Then, the iterates of LALM satisfy:

$$\|\mathbf{A}\bar{\mathbf{x}}^k - \mathbf{b}\| \leq \frac{1}{k} \left(\frac{\beta}{2} \|\mathbf{x}^0 - \mathbf{x}^*\|^2 + \frac{\max\{(1 + \|\boldsymbol{\lambda}^*\|)^2, 4\|\boldsymbol{\lambda}^*\|^2\}}{\beta} \right)$$

$$|f(\bar{\mathbf{x}}^k) - f(\mathbf{x}^*)| \leq \frac{1}{k} \left(\frac{\beta}{2} \|\mathbf{x}^0 - \mathbf{x}^*\|^2 + \frac{\max\{(1 + \|\boldsymbol{\lambda}^*\|)^2, 4\|\boldsymbol{\lambda}^*\|^2\}}{\beta} \right)$$

- Remarks:
- Guarantees are for the primal and in fact **optimal** [28].
 - No need to solve difficult subproblems at each iteration.
 - Guarantees are for $\bar{\mathbf{x}}^k$, and not \mathbf{x}^k .

Example: Basis pursuit


Problem: Basis pursuit

Given $\mathbf{A} \in \mathbb{R}^{n \times p}$ and $\mathbf{b} \in \mathbb{R}^n$, solve


$$F^* := \min_{\mathbf{x} \in \mathbb{R}^p} \{ \|\mathbf{x}\|_1 : \mathbf{Ax} = \mathbf{b} \}.$$

- ▶ Applications in de-noising, data compression.
- ▶ Experiment: \mathbf{A} is a row-normalized standard Gaussian matrix, \mathbf{x}^* is a k -sparse randomly generated vector.

Noiseless case: $\mathbf{b} := \mathbf{Ax}^*$



Noisy case: $\mathbf{b} := \mathbf{Ax}^* + \mathcal{N}(0, 10^{-3})$



Nonconvex optimization problems with nonlinear constraints

Problem template

$$f^* := \min_{\mathbf{x} \in \mathbb{R}^p} \left\{ f(\mathbf{x}) + g(\mathbf{A}(\mathbf{x})) \right\}, \quad (7)$$

- ▶ $f : \mathbb{R}^p \rightarrow \mathbb{R}$ is a proper continuously-differentiable & nonconvex
- ▶ $g : \mathbb{R}^n \rightarrow \mathbb{R}$ is proper, lower-semicontinuous
- ▶ $\mathbf{A} : \mathbb{R}^p \rightarrow \mathbb{R}^n$ is a nonlinear operator and $\mathbf{b} \in \mathbb{R}^n$
- ▶ An optimal solution \mathbf{x}^* to (7) satisfies $f(\mathbf{x}^*) = f^*$, $\mathbf{A}(\mathbf{x}^*) = \mathbf{b}$.

Example: Blind Image Deconvolution

- One of the most challenging problems in imaging sciences
 - ▶ Goal: Recover an image \mathbf{X} and an unknown blurring transformation \mathbf{T} from a blurred image $\mathbf{B} \in \mathbb{R}^{p \times q}$.
 - ▶ Formally:

$$\min_{\substack{\mathbf{T} \in \mathbb{R}^{r \times s} \\ \mathbf{X} \in \mathbb{R}^{p \times q}}} \left\{ h(\mathbf{X}, \mathbf{T}) + \frac{1}{2} \|\mathbf{T} * \mathbf{X} - \mathbf{B}\|^2 \right\},$$

where $h : \mathbb{R}^{p \times q} \times \mathbb{R}^{r \times s} \rightarrow (-\infty, +\infty]$ is a non-convex & possibly non-smooth regularizer, and $*$ is an appropriate convolution operator.

Remark: ○ Advanced material at the end of the lecture covers inexact Augmented Lagrangian for (7).

Recall the prototype problem


A primal problem prototype

$$f^* := \min_{\mathbf{x} \in \mathbb{R}^p} \left\{ f(\mathbf{x}) : \mathbf{A}\mathbf{x} = \mathbf{b}, \mathbf{x} \in \mathcal{X} \right\}, \quad (8)$$

- ▶ f is a proper, closed and convex function.
- ▶ $\mathbf{A} \in \mathbb{R}^{n \times p}$ and $\mathbf{b} \in \mathbb{R}^n$ are known.
- ▶ \mathcal{X} is nonempty, closed and convex.
- ▶ *We further assume \mathcal{X} is a bounded set! This assumption is motivated by practical applications.*

○ Standard convex optimization formulations in (8):

- ▶ *linear programming*
- ▶ *quadratic programming*
- ▶ *convex quadratic programming*
- ▶ *second order cone programming*
- ▶ *semidefinite programming*



The SDP formulation

The standard form of an SDP

$$\begin{aligned} \min_{\mathbf{X} \in \mathcal{X}} \quad & \langle \mathbf{C}, \mathbf{X} \rangle \\ \text{s.t.} \quad & \langle \mathbf{A}_i, \mathbf{X} \rangle = b_i, \text{ for } i = 1, \dots, m \end{aligned}$$

- $\mathcal{X} = \{\mathbf{X} \in \mathbb{R}^{p \times p} : \mathbf{X} \succeq 0\}$ - the positive semidefinite cone.
- $\mathbf{C} \in \mathbb{R}^{p \times p}$, $\mathbf{A}_i \in \mathbb{R}^{p \times p}$ are symmetric and $b_i \in \mathbb{R}$, and are given. By definition, $\langle \mathbf{A}_i, \mathbf{X} \rangle = \text{Tr}(\mathbf{A}_i^T \mathbf{X})$.
- Any SDP can be written in standard form.

Trace-constrained SDPs

Consider the following SDP formulation:

$$\begin{aligned} \min_{\mathbf{X} \in \mathcal{X}} \quad & \langle \mathbf{C}, \mathbf{X} \rangle \\ \text{s.t.} \quad & \langle \mathbf{A}_i, \mathbf{X} \rangle = b_i, \text{ for } i = 1, \dots, m \\ & \langle \mathbf{I}, \mathbf{X} \rangle := \text{Tr}(\mathbf{X}) = \alpha \in \mathbb{R}_+ \leftarrow \text{the trace constraint} \end{aligned} \tag{9}$$

- **Observe** that (9) belongs to the template (8).
- This formulation is of broad interest [45]. In the sequel, SDP relaxations for non-convex problems.
- Problem (9) can be large in practice, making Interior Point Methods inefficient.

Example: Finding maximum-weight cut of a graph

- **Goal:** Given an undirected graph $G = (V, E)$ with a set of weights $c : E \rightarrow \mathbb{R}_+$

$$\min_{\mathbf{x} \in \mathbb{Z}^p} \left\{ \frac{1}{2} \sum_{\{i,j\} \in E} c_{ij}(1 - x_i x_j) : x_i \in \{-1, +1\} \right\} \quad (\text{Weighted max-cut})$$

Example: Finding maximum-weight cut of a graph

- **Goal:** Given an undirected graph $G = (V, E)$ with a set of weights $c : E \rightarrow \mathbb{R}_+$

$$\min_{\mathbf{x} \in \mathbb{Z}^p} \left\{ \frac{1}{2} \sum_{\{i,j\} \in E} c_{ij}(1 - x_i x_j) : x_i \in \{-1, +1\} \right\} \quad (\text{Weighted max-cut})$$

- **The SDP approach:** Lift & relax

► lift as a matrix optimization problem $\mathbf{X} = \mathbf{x}\mathbf{x}^*$:

$$\min_{\mathbf{X} \in \mathbb{R}^{p \times p}} \left\{ \frac{1}{2} \sum_{\{i,j\} \in E} c_{ij}(1 - \mathbf{X}_{ij}) : \text{diag}(\mathbf{X}) = 1, \mathbf{X} \succeq 0, \mathbf{X}^* = \mathbf{X}, \text{rank}(\mathbf{X}) = 1 \right\}$$

► relax the non-convex rank constraint

$$\min_{\mathbf{X} \in \mathbb{R}^{p \times p}} \underbrace{\left\{ \frac{1}{2} \sum_{\{i,j\} \in E} c_{ij}(1 - \mathbf{X}_{ij}) : \underbrace{\text{diag}(\mathbf{X}) = 1}_{\mathbf{A}(\mathbf{X}) = \mathbf{b}}, \mathbf{X} \succeq 0, \mathbf{X}^* = \mathbf{X} \right\}}_{\text{tr}(\mathbf{C}\mathbf{X})} \quad (\text{Max-cut SDP})$$

Example: Finding maximum-weight cut of a graph

- **Goal:** Given an undirected graph $G = (V, E)$ with a set of weights $c : E \rightarrow \mathbb{R}_+$

$$\min_{\mathbf{x} \in \mathbb{Z}^p} \left\{ \frac{1}{2} \sum_{\{i,j\} \in E} c_{ij}(1 - x_i x_j) : x_i \in \{-1, +1\} \right\} \quad (\text{Weighted max-cut})$$

- **The SDP approach:** Lift & relax

► lift as a matrix optimization problem $\mathbf{X} = \mathbf{x}\mathbf{x}^*$:

$$\min_{\mathbf{X} \in \mathbb{R}^{p \times p}} \left\{ \frac{1}{2} \sum_{\{i,j\} \in E} c_{ij}(1 - \mathbf{X}_{ij}) : \text{diag}(\mathbf{X}) = 1, \mathbf{X} \succeq 0, \mathbf{X}^* = \mathbf{X}, \text{rank}(\mathbf{X}) = 1 \right\}$$

► relax the non-convex rank constraint

$$\min_{\mathbf{X} \in \mathbb{R}^{p \times p}} \underbrace{\left\{ \frac{1}{2} \sum_{\{i,j\} \in E} c_{ij}(1 - \mathbf{X}_{ij}) : \underbrace{\text{diag}(\mathbf{X}) = 1}_{\mathbf{A}(\mathbf{X}) = \mathbf{b}}, \mathbf{X} \succeq 0, \mathbf{X}^* = \mathbf{X} \right\}}_{\text{tr}(\mathbf{C}\mathbf{X})} \quad (\text{Max-cut SDP})$$

- Always delivers solutions 0.87856 times the optimal value after randomized rounding

Example: Clustering with minimal sum-of-squares

- **Goal:** Given data points $s_1, s_2, \dots, s_p \in \mathbb{R}^q$, assign them into k disjoint clusters.

- ▶ Minimize the sum of squared distances of all points to their cluster centers

$$\min_{\mathbf{Z}} \left\{ \sum_{j=1}^k \sum_{i=1}^p \mathbf{Z}_{ij} \|s_i - w_j(\mathbf{Z})\|^2 : \sum_{j=1}^k \mathbf{Z}_{ij} = 1, \sum_{i=1}^p \mathbf{Z}_{ij} \geq 1, \mathbf{Z}_{ij} \in \{0, 1\} \right\} \quad (\text{MinSumClu.})$$

where $\mathbf{Z} \in \{0, 1\}^{p \times k}$ is the assignment matrix with $\mathbf{Z}_{ij} = \begin{cases} 1 & \text{if } s_i \in j\text{th cluster} \\ 0 & \text{otherwise} \end{cases}$

where w_1, \dots, w_k are cluster centers with $w_j(z) = \left(\sum_{i=1}^p \mathbf{Z}_{ij} s_i \right) \left(\sum_{i=1}^p \mathbf{Z}_{ij} \right)^{-1}$

Example: Clustering with minimal sum-of-squares

- **Goal:** Given data points $s_1, s_2, \dots, s_p \in \mathbb{R}^q$, assign them into k disjoint clusters.

- ▶ Minimize the sum of squared distances of all points to their cluster centers

$$\min_{\mathbf{Z}} \left\{ \sum_{j=1}^k \sum_{i=1}^p \mathbf{Z}_{ij} \|s_i - w_j(\mathbf{Z})\|^2 : \sum_{j=1}^k \mathbf{Z}_{ij} = 1, \sum_{i=1}^p \mathbf{Z}_{ij} \geq 1, \mathbf{Z}_{ij} \in \{0, 1\} \right\} \quad (\text{MinSumClu.})$$

where $\mathbf{Z} \in \{0, 1\}^{p \times k}$ is the assignment matrix with $\mathbf{Z}_{ij} = \begin{cases} 1 & \text{if } s_i \in j\text{th cluster} \\ 0 & \text{otherwise} \end{cases}$

where w_1, \dots, w_k are cluster centers with $w_j(z) = \left(\sum_{i=1}^p \mathbf{Z}_{ij} s_i \right) \left(\sum_{i=1}^p \mathbf{Z}_{ij} \right)^{-1}$

- **The SDP approach:** Lift & relax (details omitted)

$$\min_{\mathbf{X} \in \mathbb{R}^{p \times p}} \left\{ \text{tr}(\mathbf{C}\mathbf{X}) : \mathbf{X} \geq 0, \mathbf{X}\mathbf{1} = \mathbf{1}, \mathbf{X} \succeq 0, \mathbf{X}^* = \mathbf{X}, \text{tr}(\mathbf{X}) = k \right\} \quad (\text{Clustering SDP})$$

- ▶ where $\mathbf{X} = \mathbf{Z}(\mathbf{Z}^*\mathbf{Z})^{-1}\mathbf{Z}^*$ and $c_{ij} = \|s_i - s_j\|^2$

Example: Clustering with minimal sum-of-squares

- **Goal:** Given data points $s_1, s_2, \dots, s_p \in \mathbb{R}^q$, assign them into k disjoint clusters.

- ▶ Minimize the sum of squared distances of all points to their cluster centers

$$\min_{\mathbf{Z}} \left\{ \sum_{j=1}^k \sum_{i=1}^p \mathbf{Z}_{ij} \|s_i - w_j(\mathbf{Z})\|^2 : \sum_{j=1}^k \mathbf{Z}_{ij} = 1, \sum_{i=1}^p \mathbf{Z}_{ij} \geq 1, \mathbf{Z}_{ij} \in \{0, 1\} \right\} \quad (\text{MinSumClu.})$$

where $\mathbf{Z} \in \{0, 1\}^{p \times k}$ is the assignment matrix with $\mathbf{Z}_{ij} = \begin{cases} 1 & \text{if } s_i \in j\text{th cluster} \\ 0 & \text{otherwise} \end{cases}$

where w_1, \dots, w_k are cluster centers with $w_j(z) = \left(\sum_{i=1}^p \mathbf{Z}_{ij} s_i \right) \left(\sum_{i=1}^p \mathbf{Z}_{ij} \right)^{-1}$

- **The SDP approach:** Lift & relax (details omitted)

$$\min_{\mathbf{X} \in \mathbb{R}^{p \times p}} \left\{ \text{tr}(\mathbf{C}\mathbf{X}) : \mathbf{X} \geq 0, \mathbf{X}\mathbf{1} = \mathbf{1}, \mathbf{X} \succeq 0, \mathbf{X}^* = \mathbf{X}, \text{tr}(\mathbf{X}) = k \right\} \quad (\text{Clustering SDP})$$

- ▶ where $\mathbf{X} = \mathbf{Z}(\mathbf{Z}^*\mathbf{Z})^{-1}\mathbf{Z}^*$ and $c_{ij} = \|s_i - s_j\|^2$

- Improved guarantees over LP relaxations

Example: Neural networks

- **Goal:** Approximate the ℓ_∞ -Lipschitz constant L_h of 1-layer ReLU network

$$h_{\mathbf{x}}(\mathbf{a}) := \mathbf{x}_2^T \sigma(\mathbf{X}_1 \mathbf{a} + \mathbf{x}_1)$$

- ▶ applications to verification, robustness against adversarial examples, generalization...

Example: Neural networks

- **Goal:** Approximate the ℓ_∞ -Lipschitz constant L_h of 1-layer ReLU network

$$h_{\mathbf{x}}(\mathbf{a}) := \mathbf{x}_2^T \sigma(\mathbf{X}_1 \mathbf{a} + \mathbf{x}_1)$$

- ▶ applications to verification, robustness against adversarial examples, generalization...
- **The SDP approach:** Lift & relax (details omitted)

$$L_h \leq \bar{L}_h := -\frac{1}{4} \min_{\mathbf{X} \in \mathbb{R}^{p \times p}} \{\text{tr}(\mathbf{C}\mathbf{X}) : \mathbf{X} \succeq 0, \text{diag}(\mathbf{X}) = \mathbf{1}, \mathbf{X} = \mathbf{X}^*\}$$

$$\mathbf{C} := - \begin{bmatrix} 0 & 0 & \mathbf{1}^T \mathbf{X}_2^T \text{Diag}(\mathbf{x}_2) \\ 0 & 0 & \mathbf{X}_1^T \text{Diag}(\mathbf{x}_2) \\ \text{Diag}(\mathbf{x}_2)^T \mathbf{X}_1 \mathbf{1} & \text{Diag}(\mathbf{x}_2)^T \mathbf{X}_1 & 0 \end{bmatrix}$$

Example: Neural networks

- **Goal:** Approximate the ℓ_∞ -Lipschitz constant L_h of 1-layer ReLU network

$$h_{\mathbf{x}}(\mathbf{a}) := \mathbf{x}_2^T \sigma(\mathbf{X}_1 \mathbf{a} + \mathbf{x}_1)$$

- ▶ applications to verification, robustness against adversarial examples, generalization...
- **The SDP approach:** Lift & relax (details omitted)

$$L_h \leq \bar{L}_h := -\frac{1}{4} \min_{\mathbf{X} \in \mathbb{R}^{p \times p}} \{\text{tr}(\mathbf{C}\mathbf{X}) : \mathbf{X} \succeq 0, \text{diag}(\mathbf{X}) = \mathbf{1}, \mathbf{X} = \mathbf{X}^*\}$$

$$\mathbf{C} := - \begin{bmatrix} 0 & 0 & \mathbf{1}^T \mathbf{X}_2^T \text{Diag}(\mathbf{x}_2) \\ 0 & 0 & \mathbf{X}_1^T \text{Diag}(\mathbf{x}_2) \\ \text{Diag}(\mathbf{x}_2)^T \mathbf{X}_1 \mathbf{1} & \text{Diag}(\mathbf{x}_2)^T \mathbf{X}_1 & 0 \end{bmatrix}$$

- An open research area

Ragunathan et al. SDP relaxations for certifying robustness against adversarial examples. ICLR2017

F. Latorre, P. Rolland, and V. Cevher. Lipschitz constant estimation of neural networks via sparse polynomial optimization. ICLR 2020.

CGM with quadratic penalty

Classical CGM does not apply to (3)

- lmo of the intersection of $\{x : \mathbf{Ax} = \mathbf{b}\}$ and \mathcal{X} is difficult to compute.
- Idea: Combine the CGM framework with the quadratic penalty approach.

Quadratic penalty strategy

- A quadratic penalty formulation:

$$\min_{\mathbf{x} \in \mathbb{R}^p} \left\{ f(\mathbf{x}) + \frac{\beta}{2} \|\mathbf{Ax} - \mathbf{b}\|_2^2 : \mathbf{x} \in \mathcal{X} \right\}$$

- $\beta > 0$ is the penalty parameter and $f_\beta(\mathbf{x})$ is the penalized objective function.
 - Note that $f_\beta(\mathbf{x})$ is convex and smooth with parameter $L + \beta \|\mathbf{A}\|^2$.
- A simple strategy [42] \Rightarrow Take a CGM step on f_β and increase β progressively

Homotopy conditional gradient method (HCGM)

1. Choose $\mathbf{x}^0 \in \mathcal{X}$, and $\beta_0 > 0$.

2. For $k = 0, 1, \dots$:

$$\hat{\mathbf{x}}^k := \text{lmo}_{\mathcal{X}}(\nabla f(\mathbf{x}^k) + \beta_k \mathbf{A}^T(\mathbf{Ax}^k - \mathbf{b})).$$

$$\mathbf{x}^{k+1} := (1 - \gamma_k) \mathbf{x}^k + \gamma_k \hat{\mathbf{x}}^k,$$

where $\gamma_k := \frac{2}{k+2}$ and $\beta_k := \beta_0 \sqrt{k+2}$.

Convergence guarantees of HCGM

Recall Lagrange duality

$$\mathcal{L}(\mathbf{x}, \boldsymbol{\lambda}) := f(\mathbf{x}) + \langle \boldsymbol{\lambda}, \mathbf{A}\mathbf{x} - \mathbf{b} \rangle$$
$$\underbrace{\max_{\boldsymbol{\lambda}} \min_{\mathbf{x} \in \mathcal{X}} \mathcal{L}(\mathbf{x}, \boldsymbol{\lambda})}_{\text{dual problem}} \leq \underbrace{\min_{\mathbf{x} \in \mathcal{X}} \max_{\boldsymbol{\lambda}} \mathcal{L}(\mathbf{x}, \boldsymbol{\lambda})}_{\text{primal problem}} \quad (\text{Duality})$$

- $\boldsymbol{\lambda}$ is called the **Lagrange multiplier**.
- The function $d(\boldsymbol{\lambda})$ is called the **dual function**, and it is **concave**!
- The optimal dual objective value is $d^* = d(\boldsymbol{\lambda}^*)$.

(Duality) holds with equality under weak assumptions \Rightarrow (Strong duality).

Theorem (Simplified[42])

Assume that strong duality holds. Then, the iterates of HCGM satisfy

$$\begin{cases} |f(\mathbf{x}^k) - f^*| & \in \mathcal{O}(k^{-1/2}) \\ \|\mathbf{A}\mathbf{x}^k - \mathbf{b}\| & \in \mathcal{O}(k^{-1/2}). \end{cases}$$

* For an extension of HCGM to the case $\mathbf{A}\mathbf{x} - \mathbf{b} \in \mathcal{K}$, please see Appendix A₁.

** Advanced material at the end of the lecture covers stochastic variants of HCGM.

Augmented Lagrangian CGM: CGAL

Augmented Lagrangian approach

- Augmented problem formulation:

$$\min_{\mathbf{x} \in \mathbb{R}^p} \left\{ f(\mathbf{x}) + \frac{\beta}{2} \|\mathbf{Ax} - \mathbf{b}\|_2^2 : \mathbf{Ax} = \mathbf{b}, \mathbf{x} \in \mathcal{X} \right\}$$

- Write down the Lagrangian:

$$\mathcal{L}_\beta(\mathbf{x}, \boldsymbol{\lambda}) = f(\mathbf{x}) + \langle \boldsymbol{\lambda}, \mathbf{Ax} - \mathbf{b} \rangle + \frac{\beta}{2} \|\mathbf{Ax} - \mathbf{b}\|^2$$

- Note that $\mathcal{L}_\beta(\cdot, \boldsymbol{\lambda})$ is smooth with parameter $L + \beta \|\mathbf{A}\|^2$.

- Our strategy [40] \Rightarrow
 - Take a CGM step wrt $\mathcal{L}_\beta(\cdot, \boldsymbol{\lambda})$ (primal)
 - Take a gradient step wrt $\mathcal{L}_\beta(\mathbf{x}, \cdot)$ (dual)
 - Increase β progressively
- Challenge: Step size in the dual domain (step 2.)

Convergence guarantees of CGAL

| Conditional gradient augmented Lagrangian method (CGAL) | |
|---|--|
| 1. | Choose $\mathbf{x}^0 \in \mathcal{X}$, $\boldsymbol{\lambda}^0 \in \mathbb{R}^n$, and $\beta_0 > 0$. |
| 2. | For $k = 0, 1, \dots$: |
| | $\hat{\mathbf{x}}^k := \text{lmo}_{\mathcal{X}}(\nabla f(\mathbf{x}^k) + \mathbf{A}^T \boldsymbol{\lambda}^k + \beta_k \mathbf{A}^T (\mathbf{Ax}^k - \mathbf{b}))$ |
| | $\mathbf{x}^{k+1} := (1 - \gamma_k) \mathbf{x}^k + \gamma_k \hat{\mathbf{x}}^k$ |
| | $\boldsymbol{\lambda}^{k+1} := \boldsymbol{\lambda}^k + \omega_k (\mathbf{Ax}^{k+1} - \mathbf{b})$ |
| | where $\gamma_k := \frac{2}{k+2}$, and $\beta_k := \beta_0 \sqrt{k+2}$. |

Theorem (Simplified)

Assume that strong duality holds. Let us choose dual step size ω_k by the following rule

$$\omega_k = \alpha_k := \min \left\{ \frac{1}{\beta_0}, \frac{\eta_k^2 (L_f + \boldsymbol{\lambda}_{k+1}) D_{\mathcal{X}}^2}{2 \|\mathbf{Ax}^{k+1} - \mathbf{b}\|^2} \right\} \quad \text{if} \quad \|\boldsymbol{\lambda}^k + \alpha_k (\mathbf{Ax}^{k+1} - \mathbf{b})\| \leq D_{\mathcal{Y}}$$



and $\omega_k = 0$ otherwise, for some $D_{\mathcal{Y}} \geq 0$. Then, the iterates of CGAL satisfy

$$\begin{cases} |f(\mathbf{x}^k) - f^*| & \in \mathcal{O}(\frac{1}{\sqrt{k}}) \\ \|\mathbf{Ax}^k - \mathbf{b}\| & \in \mathcal{O}(\frac{1}{\sqrt{k}}) \end{cases}$$

* For an extension of CGAL to the case $\mathbf{Ax} - \mathbf{b} \in \mathcal{K}$, please see Appendix A₂.

Example: k-means clustering


$$\min_{\mathbf{X} \in \mathbb{R}^{p \times p}} \left\{ \text{Tr}(\mathbf{C}\mathbf{X}) : \mathbf{X}\mathbf{1} = \mathbf{1}, \mathbf{X} \geq 0, \mathbf{X} \in \mathcal{S}_+^p, \text{Tr}(\mathbf{X}) = \alpha \right\}$$



- ▶ Test setup with preprocessed MNIST dataset [42]
- ▶ $p = 1000$ & $\alpha = 10$ is the number of clusters
- ▶ Note: the worst-case guarantee is the same for HCGM and CGAL, but CGAL performs better in practice.

Example: Max-cut SDP

$$\max_{\mathbf{X} \in \mathbb{R}^{p \times p}} \left\{ \frac{1}{4} \text{Tr}(\mathbf{L}\mathbf{X}) : \text{diag}(\mathbf{X}) = \mathbf{1}, \mathbf{X} \in \mathcal{S}_+^p, \text{Tr}(\mathbf{X}) = p \right\}$$



- UF Sparse graphs: GSet collection, G40 dataset $p = 2000$
- \mathbf{L} is graph Laplacian matrix.
- Note: the worst-case guarantee is the same for HCGM and CGAL, but CGAL performs better **in practice**.

Towards scalable semidefinite programming

Structures in SDP relaxations

$$\min_{\mathbf{X} \in \mathbb{R}^{p \times p}} \{\text{Tr}(\mathbf{C}\mathbf{X}) : \mathcal{A}\mathbf{X} = b, \mathbf{X} \succeq 0, \text{Tr}(\mathbf{X}) = \alpha\} \quad (10)$$

- \mathbf{X} has $\mathcal{O}(p^2)$ -degrees of freedom \implies needs $\Theta(p^2)$ storage
- Optimal solutions \mathbf{X}^* typically or approximately have $\mathcal{O}(rp)$ -degrees of freedom
 - ▶ $r = \text{rank}$ & $r \ll p$ (*low-rank*)
 - ▶ \implies need $\Theta(rp)$ storage for a rank- r approximate solution
- Example SDP's typically have $n = \tilde{\mathcal{O}}(p)$ affine constraints
 - ▶ During optimization we need to keep track of quantities such as

$$A(uv^*) - u^*(A^*z) - (A^*z)v, \quad u \in \mathbb{R}^p, v \in \mathbb{R}^p, z \in \mathbb{R}^n$$

\implies need $\Omega(n + p)$ storage for computations

Towards scalable semidefinite programming

Structures in SDP relaxations

$$\min_{\mathbf{X} \in \mathbb{R}^{p \times p}} \{\text{Tr}(\mathbf{C}\mathbf{X}) : \mathcal{A}\mathbf{X} = b, \mathbf{X} \succeq 0, \text{Tr}(\mathbf{X}) = \alpha\} \quad (10)$$

- \mathbf{X} has $\mathcal{O}(p^2)$ -degrees of freedom \implies needs $\Theta(p^2)$ storage ← this becomes a major problem
- Optimal solutions \mathbf{X}^* typically or approximately have $\mathcal{O}(rp)$ -degrees of freedom
 - ▶ $r = \text{rank } \mathbf{X}^*$ & $r \ll p$ (*low-rank*)
 - ▶ \implies need $\Theta(rp)$ storage for a rank- r approximate solution
- Example SDP's typically have $n = \tilde{\mathcal{O}}(p)$ affine constraints
 - ▶ During optimization we need to keep track of quantities such as

$$A(uv^*) \quad u^*(A^*z) \quad (A^*z)v, \quad u \in \mathbb{R}^p, v \in \mathbb{R}^p, z \in \mathbb{R}^n$$

\implies need $\Omega(n + p)$ storage for computations

$\Theta(n + rp)$ storage

Towards scalable semidefinite programming

Structures in SDP relaxations

$$\min_{\mathbf{X} \in \mathbb{R}^{p \times p}} \{\text{Tr}(\mathbf{C}\mathbf{X}) : \mathcal{A}\mathbf{X} = b, \mathbf{X} \succeq 0, \text{Tr}(\mathbf{X}) = \alpha\} \quad (10)$$

- \mathbf{X} has $\mathcal{O}(p^2)$ -degrees of freedom \implies needs $\Theta(p^2)$ storage ← this becomes a major problem

- Optimal solutions \mathbf{X}^* typically or approximately have $\mathcal{O}(rp)$ -degrees of freedom

- ▶ $r = \text{rank}$ & $r \ll p$ (*low-rank*)
 - ▶ \implies need $\Theta(rp)$ storage for a rank- r approximate solution

- Example SDP's typically have $n = \tilde{\mathcal{O}}(p)$ affine constraints

- ▶ During optimization we need to keep track of quantities such as

$$A(uv^*) \quad u^*(A^*z) \quad (A^*z)v, \quad u \in \mathbb{R}^p, v \in \mathbb{R}^p, z \in \mathbb{R}^n$$

\implies need $\Omega(n + p)$ storage for computations

$\Theta(n + rp)$ storage

- ▶ Relevant SDPs are often large \implies HCGM, CGAL have a storage bottleneck (e.g., MaxCut for graph of $2e^6$ nodes $\rightarrow \sim 2e^{12}$ variables!!)
- ▶ Can we leverage the **problem structure** for better storage performance? See advanced material.

Wrap up!

- Last lecture! HW and mock exam...

*An explicit inexact ALM: ASGARD-DL

Inexact ALM (Double Loop ASGARD [35])

1. $\mathbf{x}^0 = \hat{\mathbf{x}}^{0,0} = \bar{\mathbf{x}}^{0,0} = \tilde{\mathbf{x}}^{0,0} \in \mathbb{R}^p$, $\boldsymbol{\lambda}_0 \in \mathbb{R}^n$, $\beta_k > 0$, $\tau_0 = 1$, $m_0 > 2$, $\omega > 1$.

2. For $k = 0, 1, \dots$, perform:

2.a For $i = 0, 1, \dots, m_k - 1$: // accelerated proximal method

$$\hat{\mathbf{x}}^{k,i} = (1 - \tau_k) \bar{\mathbf{x}}^{k,i} + \tau_k \tilde{\mathbf{x}}^{k,i}$$

$$\tilde{\mathbf{x}}^{k,i+1} = \text{prox}_{\frac{f}{\beta_k \|A\|^2}} \left(\tilde{\mathbf{x}}^{k,i} - \frac{1}{\beta_k \|A\|^2} A^\top (\boldsymbol{\lambda}^k + \beta_k (A \hat{\mathbf{x}}^{k,i} - \mathbf{b})) \right)$$

$$\bar{\mathbf{x}}^{k,i+1} = \hat{\mathbf{x}}^{k,i} + \tau_k (\tilde{\mathbf{x}}^{k,i+1} - \tilde{\mathbf{x}}^{k,i})$$

$$\tau_{k+1} = \frac{2}{k+2}$$

2.b Update primal and dual variables:

$$\bar{\mathbf{x}}^{k+1,0} = \tilde{\mathbf{x}}^{k,m_k}$$

$$\boldsymbol{\lambda}^{k+1} = \boldsymbol{\lambda}^k + \beta_k (A \bar{\mathbf{x}}^{k+1,0} - \mathbf{b}), \quad // \text{ update dual variable}$$

$$\tau_0 = 1$$

$$\beta_{k+1} = \beta_k \omega, \quad // \text{ increase } \beta_k$$

$$m_{k+1} = m_k \omega, \quad // \text{increase # of inner iterations}$$

Remarks:

- Corresponds to inexact ALM with explicit inner termination rule.
- Attains optimal $\mathcal{O}(1/k)$ on the last iterate, with good empirical performance (see slide 17).

*ADMM¹

Primal problem with a specific decomposition structure

$$f^* := \min_{\mathbf{x} := (\mathbf{u}, \mathbf{v})} \{ f(\mathbf{x}) := g(\mathbf{u}) + h(\mathbf{v}) : \mathbf{B}\mathbf{u} + \mathbf{C}\mathbf{v} = \mathbf{b}, \mathbf{u} \in \mathcal{U}, \mathbf{v} \in \mathcal{V} \}$$

- $\mathcal{X} := \mathcal{U} \times \mathcal{V}$ - nonempty, closed, convex and **bounded**.
- $\mathbf{A} := [\mathbf{B}, \mathbf{C}]$.

The Fenchel dual problem

$$d^* := \max_{\lambda \in \mathbb{R}^n} \left\{ d(\lambda) := -g_{\mathcal{U}}^*(-\mathbf{B}^T \lambda) - h_{\mathcal{V}}^*(-\mathbf{C}^T \lambda) + \langle \mathbf{b}, \lambda \rangle \right\}$$

- $g_{\mathcal{U}}^*$ and $h_{\mathcal{V}}^*$ are the Fenchel conjugate of $g_{\mathcal{U}} := g + \delta_{\mathcal{U}}$ and $h_{\mathcal{V}} := h + \delta_{\mathcal{V}}$, resp.

The dual function

$$d(\lambda) := \underbrace{\min_{\mathbf{u} \in \mathcal{U}} \left\{ g(\mathbf{u}) + \langle \mathbf{B}^T \lambda, \mathbf{u} \rangle \right\}}_{d^1(\lambda)} + \underbrace{\min_{\mathbf{v} \in \mathcal{V}} \left\{ h(\mathbf{v}) + \langle \mathbf{C}^T \lambda, \mathbf{v} \rangle \right\}}_{d^2(\lambda)} - \langle \mathbf{b}, \lambda \rangle.$$

*Splitting the problem

Standard ADMM


$$\begin{cases} \mathbf{u}^{k+1} &:= \arg \min_{\mathbf{u} \in \mathcal{U}} \left\{ g(\mathbf{u}) + \langle \lambda^k, \mathbf{B}\mathbf{u} \rangle + \frac{\beta_k}{2} \|\mathbf{B}\mathbf{u} + \mathbf{C}\mathbf{v}^k - \mathbf{b}\|^2 \right\} \\ \mathbf{v}^{k+1} &:= \arg \min_{\mathbf{v} \in \mathcal{V}} \left\{ h(\mathbf{v}) + \langle \lambda^k, \mathbf{C}\mathbf{v} \rangle + \frac{\beta_k}{2} \|\mathbf{B}\mathbf{u}^{k+1} + \mathbf{C}\mathbf{v} - \mathbf{b}\|^2 \right\} \\ \lambda^{k+1} &:= \lambda^k + \beta_k (\mathbf{B}\mathbf{u}^{k+1} + \mathbf{C}\mathbf{v}^{k+1} - \mathbf{b}). \end{cases}$$

Here, $\beta_k > 0$ is a given **penalty parameter** of the associated augmented problem:

$$\mathcal{L}_\beta := g(\mathbf{u}) + h(\mathbf{v}) + \langle \lambda, \mathbf{B}\mathbf{u} + \mathbf{C}\mathbf{v} - \mathbf{b} \rangle + \frac{\beta}{2} \|\mathbf{B}\mathbf{u} + \mathbf{C}\mathbf{v} - \mathbf{b}\|^2$$

- Note how minimizing over (\mathbf{u}, \mathbf{v}) together would reduce to the ALM formulation.

Leveraging GANs for Signal Recovery



Problem formulation

$$\min_{\mathbf{w}, \mathbf{z}} L(\mathbf{w}) + R(\mathbf{w}) + H(\mathbf{z}) \quad \text{subject to } \mathbf{w} = G(\mathbf{z})$$

- ▶ L is convex and smooth
- ▶ R, H are convex and proximal friendly
- ▶ G differentiable generative model (non-linear and usually non-convex)

*AL schemes for non-convex problems - challenges

Challenges

- More complicated requirements to prove global convergence of generic schemes for (7) (e.g., [31]):
 - ▶ \exists superset of the feasible-set, where feasibility is ‘good-enough’ (information zone - IZ)
 - ▶ Objective & constraints need to be ‘sufficiently-regular’ within the IZ
 - ▶ The iterates of the AL algorithm need to
 - ▶ Enter the IZ in a finite number of steps.
 - ▶ Stay inside the IZ thereafter.
- Literature studying this setting is scarce, and global convergence is not well-understood.
- A practically-relevant variation of (7) has recently been analyzed via the inexact AL scheme [32]. ← up next

Sam

Set-up

Assume the following template:

$$\min_{\mathbf{x} \in \mathbb{R}^p} f(\mathbf{x}) + g(\mathbf{x}) \text{ s.t. } \mathbf{A}(\mathbf{x}) = \mathbf{b} \quad (11)$$

- ▶ $f : \mathbb{R}^p \rightarrow \mathbb{R}$ is a continuously-differentiable non-convex function that is L_f -smooth.
- ▶ $g : \mathbb{R}^p \rightarrow \mathbb{R}$ is a proximal-friendly convex function.
- ▶ $\mathbf{A} : \mathbb{R}^p \rightarrow \mathbb{R}^n$ is a smooth nonlinear operator i.e., $\exists L_{\mathbf{A}} > 0$ s.t.: $\|\mathbf{J}_{\mathbf{A}}(\mathbf{x}) - \mathbf{J}_{\mathbf{A}}(\mathbf{y})\| \leq L_{\mathbf{A}} \|\mathbf{x} - \mathbf{y}\|$, where \mathbf{J} is the Jacobian of \mathbf{A} .

*AL schemes for non-convex problems - challenges

Challenges

- More complicated requirements to prove global convergence of generic schemes for (7) (e.g., [31]):
 - ▶ \exists superset of the feasible-set, where feasibility is ‘good-enough’ (information zone - IZ)
 - ▶ Objective & constraints need to be ‘sufficiently-regular’ within the IZ
 - ▶ The iterates of the AL algorithm need to
 - ▶ Enter the IZ in a finite number of steps.
 - ▶ Stay inside the IZ thereafter.
- Literature studying this setting is scarce, and global convergence is not well-understood.
- A practically-relevant variation of (7) has recently been analyzed via the inexact AL scheme [32]. ← up next

Sam

Set-up

Assume the following template:

$$\min_{\mathbf{x} \in \mathbb{R}^p} f(\mathbf{x}) + g(\mathbf{x}) \text{ s.t. } \mathbf{A}(\mathbf{x}) = \mathbf{b} \quad (12)$$

- ▶ $f : \mathbb{R}^p \rightarrow \mathbb{R}$ is a continuously-differentiable non-convex function that is L_f -smooth.
- ▶ $g : \mathbb{R}^p \rightarrow \mathbb{R}$ is a proximal-friendly convex function.
- ▶ $\mathbf{A} : \mathbb{R}^p \rightarrow \mathbb{R}^n$ is a smooth nonlinear operator i.e., $\exists L_{\mathbf{A}} > 0$ s.t.: $\|\mathbf{J}_{\mathbf{A}}(\mathbf{x}) - \mathbf{J}_{\mathbf{A}}(\mathbf{y})\| \leq L_{\mathbf{A}} \|\mathbf{x} - \mathbf{y}\|$, where \mathbf{J} is the Jacobian of \mathbf{A} .

*AL schemes for non-convex problems - optimality conditions

Reformulating (12) in terms of AL

- Solving (12) is equivalent to solving the following reformulation:

$$\min_{\mathbf{x}} \max_{\boldsymbol{\lambda}} \mathcal{L}_\beta(\mathbf{x}, \boldsymbol{\lambda}) + g(\mathbf{x})$$

where for a given $\beta > 0$, $\mathcal{L}_\beta(\mathbf{x}, \boldsymbol{\lambda}) = f(\mathbf{x}) + \langle \mathbf{A}(\mathbf{x}) - \mathbf{b}, \boldsymbol{\lambda} \rangle + \frac{\beta}{2} \|\mathbf{A}(\mathbf{x}) - \mathbf{b}\|^2$ - the Augmented Lagrangian.

Optimality conditions of (12)

- $\mathbf{x} \in \mathbb{R}^p$ is a first order stationary point (FOS) of (12) if $\exists \boldsymbol{\lambda} \in \mathbb{R}^n$ s.t.

$$-\nabla \mathcal{L}_\beta(\mathbf{x}, \boldsymbol{\lambda}) \in \partial g(\mathbf{x}) \quad \text{and} \quad \mathbf{A}(\mathbf{x}) = \mathbf{b}.$$

- When $g = 0$ and \mathbf{x} is a FOS, \mathbf{x} is also a second-order stationary point (SOS) if:

$$\lambda_{\min} (\nabla_{\mathbf{x}\mathbf{x}} \mathcal{L}_\beta(\mathbf{x}, \boldsymbol{\lambda})) \geq 0$$

- Approximate stationarity is then defined for a given $\epsilon > 0$ as:

- ▶ FOS: $\text{dist}(-\nabla \mathcal{L}_\beta(\mathbf{x}, \boldsymbol{\lambda}), \partial g(\mathbf{x})) \leq \epsilon \quad \text{and} \quad \|\mathbf{A}(\mathbf{x}) - \mathbf{b}\| \leq \epsilon$

- ▶ SOS: $\lambda_{\min} (\nabla_{\mathbf{x}\mathbf{x}} \mathcal{L}_\beta(\mathbf{x}, \boldsymbol{\lambda})) \geq -\epsilon$

*An Inexact AL scheme for non-convex problems

- Main idea of [32]: solve primal problems with increasing accuracy ϵ_k and carefully choose the dual stepsize σ_k .

| ALM - conceptual (reference) | Inexact ALM - nonconvex (IALM) |
|---|---|
| <p>1. Choose $\lambda_0 \in \mathbb{R}^n$ and $\beta > 0$.</p> <p>2. For $k = 0, 1, \dots$:</p> <p>2.a. Solve (4) to get \mathbf{x}^{k+1}.</p> <p>2.b. Update $\lambda^{k+1} := \lambda^k + \beta \left(\mathbf{A}\mathbf{x}_\beta^*(\lambda^k) - \mathbf{b} \right)$.</p> | <p>1. Choose $b > 1$, $\lambda^0 \in \mathbb{R}^n$, $\sigma_0 > 0$, τ_f, $\tau_s > 0$.</p> <p>2. For $k = 0, 1, \dots$, perform:</p> <p>2.aa. Set $\epsilon_{k+1} = 1/\beta_k$</p> <p>2.a. Get \mathbf{x}^{k+1} with a solver of choice, s.t.:</p> $\text{dist}(-\nabla_x \mathcal{L}_{\beta_k}(\mathbf{x}^{k+1}, \lambda_k), \partial g(\mathbf{x}^{k+1})) \leq \epsilon_{k+1}, \quad [\text{FOS}]$ <p style="text-align: center;">or</p> $\lambda_{\min}(\nabla_{\mathbf{x}\mathbf{x}} \mathcal{L}_{\beta_k}(\mathbf{x}^{k+1}, \lambda^k)) \geq -\epsilon_{k+1} \quad [\text{SOS}]$ <p>2.b. Update</p> $\beta_{k+1} = b^{k+1}$ $\sigma_{k+1} = \sigma_0 \min \left(1, \frac{\ \mathbf{A}(\mathbf{x}^1) - \mathbf{b}\ \log^2(2)}{\ \mathbf{A}(\mathbf{x}^{k+1}) - \mathbf{b}\ (k+1) \log^2(k+2)} \right)$ $\lambda^{k+1} = \lambda^k + \sigma_{k+1} (\mathbf{A}(\mathbf{x}^{k+1}) - \mathbf{b})$ <p>2.c. Stop if</p> $\text{dist}(-\nabla_x \mathcal{L}_{\beta_k}(\mathbf{x}^{k+1}, \lambda_k), \partial g(\mathbf{x}^{k+1})) + \ \mathbf{A}(\mathbf{x}^{k+1}) - \mathbf{b}\ \leq \tau_f \quad [\text{FOS}]$ <p>and if also $\lambda_{\min}(\nabla_{\mathbf{x}\mathbf{x}} \mathcal{L}_{\beta_k}(\mathbf{x}^{k+1}, \lambda^k)) \geq -\epsilon_{k+1} \quad [\text{SOS}]$</p> |

*Convergence of the Inexact AL for non-convex problems

A key assumption

- For convex AL schemes we rely on Slater's condition to prove convergence.
- We need a similar kind of assumption for our non-convex problem, called **regularity condition**²: for some $\nu > 0$, assume

$$\nu \|\mathbf{A}(\mathbf{x}^k) - \mathbf{b}\| \leq \text{dist} \left(-\mathbf{J}_{\mathbf{A}}(\mathbf{x}^k)^{\top} (\mathbf{A}(\mathbf{x}^k) - \mathbf{b}), \frac{\partial g(\mathbf{x}^k)}{\beta_{k-1}} \right), \quad \forall k \quad (13)$$

- Informally, condition (13) ensures that step 2.a of IALM improves feasibility as β_k grows.

Theorem [32] (Simplified)

Under the framework (12) and assumption (13), IALM reaches

- ▶ FOS with $\tilde{\mathcal{O}}(\epsilon^3)$ complexity and
- ▶ SOS with $\tilde{\mathcal{O}}(\epsilon^5)$ complexity,

where $\tilde{\mathcal{O}}$ hides logarithmic factors³.


*Example: k-means clustering

- Model free k-means clustering SDP:

$$\min_{\mathbf{X} \in \mathbb{R}^{p \times p}} \left\{ \text{Tr}(\mathbf{C}\mathbf{X}) : \mathbf{X}\mathbf{1} = \mathbf{1}, \mathbf{X} \geq 0, \mathbf{X} \in \mathcal{S}_+^p, \text{Tr}(\mathbf{X}) = \alpha \right\}$$

- Nonconvex formulation:

$$\min_{\mathbf{u} \in \mathbb{R}^p} \left\{ \text{Tr}(\mathbf{C}\mathbf{u}\mathbf{u}^*) : \mathbf{u}\mathbf{u}^*\mathbf{1} = \mathbf{1}, \mathbf{u} \geq 0, \|\mathbf{u}\|_F \leq \sqrt{\alpha} \right\},$$



*Example: DARN with GANs (MNIST)

- De-adversarial-noise with generative adversarial networks:

$$\min_{\mathbf{w}, \mathbf{z}} \{ \| \mathbf{w} - (\mathbf{w}_0 + \eta) \|_* : \mathbf{w} = \mathbf{G}(\mathbf{z}) \}$$





Figure: ℓ_∞ error per iteration

Mathematics of Data | Prof. Volkan Cevher, volkan.cevher@epfl.ch

Figure: misclassification error per iteration

*Example: Basis Pursuit



- Convex formulation:

$$\min_{\mathbf{x} \in \mathbb{R}^p} \left\{ \|\mathbf{x}\|_1 : \mathbf{Ax} = \mathbf{b} \right\}$$

- Non-convex formulation:

change of variables

$$\begin{cases} \mathbf{x} &:= \mathbf{x}^+ - \mathbf{x}^- \\ \mathbf{x}^+ &:= \mathbf{u}_1^{\circ 2}, \quad \mathbf{x}^- := \mathbf{u}_2^{\circ 2} \text{ and } \mathbf{u} := [\mathbf{u}_1^\top, \mathbf{u}_2^\top]^\top \\ \bar{\mathbf{A}} &:= [\mathbf{A}, -\mathbf{A}] \end{cases} \rightarrow \min_{\mathbf{u} \in \mathbb{R}^p} \left\{ \|\mathbf{u}\|_2^2 : \bar{\mathbf{A}}\mathbf{u}^{\circ 2} = \mathbf{b} \right\}$$



*Stochastic HCGM for almost sure constraints

Problem formulation


$$f^* := \min_{x \in \mathcal{X}} f(x) := \mathbb{E}[f(x, \xi)], \quad A(\xi)x \in b(\xi) \text{ a. s.,}$$

- ▶ $f(x, \xi) : \mathbb{R}^d \rightarrow \mathbb{R}$ - convex, L_f -Lipschitz gradient
- ▶ $\mathcal{X} \subset \mathbb{R}^d$ - convex, compact
- ▶ $A(\xi) \in \mathbb{R}^{m \times d}$ - matrix-valued random variable, $b(\xi) \subset \mathbb{R}^m$ - random convex set

Applications

- ▶ Target application: solving large scale SDPs
- ▶ Stochastic template \implies formulation of stochastic first order methods \implies can handle large problems, in both dimension and constraints
- ▶ Two examples relevant to ML:
 - ▶ K-Means clustering SDP [29, 24]
 - ▶ Sparsest Cut SDP [1, 8]

*Main idea



- ▶ Optimize increasingly accurate approximations $F_\beta(x)$, as $\beta \rightarrow 0$
- ▶ Control the exploding variance using variance reduction

*First method: H-1SFW

Algorithm - H-1SFW($x_1 \in \mathcal{X}, \beta_0 > 0, P(\xi)$)

for $k = 1, 2, \dots$ **do**

 Set $\rho_k, \gamma_k, \beta_k$, sample $\xi_k \sim P(\xi) \leftarrow \gamma_k \in \mathcal{O}(1/k), \beta_k \in \mathcal{O}(1/\sqrt{k})$

$v_k = (1 - \rho_k)v_{k-1} + \rho_k \nabla_x F_{\beta_k}(x_k, \xi_k) \leftarrow$ variance reduction on gradient estimator v_k with single sample ξ_k [25]

$x_{k+1} = \text{fw_step}(x_k, v_k, \gamma_k) \leftarrow$ the usual lmo $_{\mathcal{X}}(v_k)$ and convex combination update

end for

Convergence H-1SFW

If $\mathbb{E}[\nabla f(x, \xi)] = \nabla f(x)$, $\mathbb{E}[\|\nabla f(x, \xi) - \nabla f(x)\|^2] \leq \sigma_f^2 < +\infty$, $\sup_{\xi} \|A(\xi)\|^2 < +\infty$ and Slater's condition holds, then for all k :

$$\mathbb{E}[|f(x_k, \xi) - f(x_*)|] \in \mathcal{O}(k^{-1/6}), \quad \sqrt{\mathbb{E}[\text{dist}(A(\xi)x_k, b(\xi))^2]} \in \mathcal{O}(k^{-1/6})$$

The oracle complexity for ϵ -accuracy is: $\mathcal{O}(\epsilon^{-6})$ stochastic first order oracles (#sfo) and $\mathcal{O}(\epsilon^{-6})$ linear minimization oracles (#lmo).

*Second method: H-SPIDER-FW

Algorithm - H-SPIDER-FW($\bar{x}_1 \in \mathcal{X}, \beta_0 > 0, P(\xi)$)

for $t = 1, 2, \dots$ **do**

 Set $x_{t,1} = \bar{x}_t; K_t = 2^t; \gamma_{t,1}; \beta_{t,1};$, sample $\xi_{\mathcal{Q}_t} \stackrel{i.i.d.}{\sim} P(\xi)$ \leftarrow Set minibatch size K_t

$v_{t,1} = \tilde{\nabla} F_{\beta_{t,1}}(x_{t,1}, \xi_{\mathcal{Q}_t})$ \leftarrow Compute 'high-accuracy' averaged stochastic gradient

$x_{t,2} = \text{fw_step}(x_{t,1}, v_{t,1}, \gamma_{t,1})$

for $k = 2, \dots, K_t$ **do**

 Set $\gamma_{t,k}; \beta_{t,k}$, sample $\xi_{\mathcal{S}_{t,k}} \stackrel{i.i.d.}{\sim} P(\xi)$ \leftarrow Decrease $\beta \in \mathcal{O}\left(1/\sqrt{K_t + k}\right)$, set $\gamma \in \mathcal{O}(1/(K_t + k))$

$v_{t,k} = v_{t,k-1} - \tilde{\nabla} F_{\beta_{t,k-1}}(x_{t,k-1}, \xi_{\mathcal{S}_{t,k}}) + \tilde{\nabla} F_{\beta_{t,k}}(x_{t,k}, \xi_{\mathcal{S}_{t,k}})$ \leftarrow var. red. on v_k using minibatch [43]

$x_{t,k+1} = \text{fw_step}(x_{t,k}, v_{t,k}, \gamma_{t,k})$ \leftarrow the usual $\text{lmo}_{\mathcal{X}}(v_k)$ and convex combination update

end for

$\bar{x}_{t+1} = x_{t,K_t+1}$

end for

Convergence H-SPIDER-FW

Denote by $n := K_t + k$ the global iteration #. Under identical assumptions as H-1SFW, for all k it holds that:


$$\mathbb{E} [|f(x_{t,k}, \xi) - f(x_*)|] \in \mathcal{O}(n^{-1/2}), \quad \sqrt{\mathbb{E} [\text{dist}(A(\xi)x_{t,k}, b(\xi))^2]} \in \mathcal{O}(n^{-1/2})$$

The oracle complexity for ϵ -accuracy is: $\mathcal{O}(\epsilon^{-2}) \#sfo$ and $\mathcal{O}(\epsilon^{-4}) \#\text{lmo}$.

*Experiments: Uniform Sparsest Cut SDP

- Approximation algorithm [1] is based on SDP relaxation: dimension $\mathcal{O}(d^2)$, constraints $\mathcal{O}(d^3)$

$$\begin{aligned} & \min_{X \in \mathbf{X}} \quad \langle L, X \rangle \\ \text{subject to} \quad & d\text{Tr}(X) - \text{Tr}(\mathbf{1}_{d \times d} X) = \frac{d^2}{2} \\ & X_{i,j} + X_{j,k} - X_{i,k} - X_{j,j} \leq 0 \quad \forall i, j, k \in V \end{aligned}$$



- From left to right (columns): 25 nodes, $\sim 7e^3$ constraints; 55 nodes, $\sim 8e^4$ constraints; 102 nodes, $\sim 5e^5$ constraints. Graphs from [?]

*Towards scalable semidefinite programming

The road to storage optimality

- ▶ SDPs often have a low rank solutions \implies instead of storing $\mathbf{X}_{k \in \{1 \dots T\}}$ at every iteration, use a compressed representation S_k given by a [matrix sketching](#) technique.
- ▶ [Formally](#) - Consider a PSD matrix $\mathbf{X} \in \mathbb{R}^{p \times p}$ and let $R > 0$ be a parameter that controls the storage cost of a sketch (and its accuracy). Construct a so-called Nyström sketch by drawing a fixed standard normal matrix $\Omega \in \mathbb{R}^{p \times R}$, and produce a sketch \mathbf{S} of \mathbf{X} as follows:

$$\mathbf{S} = \mathbf{X}\Omega \in \mathbb{R}^{p \times R}$$

- ▶ [Reconstruction](#) - Given Ω and \mathbf{S} , we recover a rank-R approximation $\hat{\mathbf{X}}$ of \mathbf{X} by

$$\hat{\mathbf{X}} := \mathbf{S}(\Omega^T \mathbf{S})^\dagger \mathbf{S}^T \quad \text{with} \quad \mathbb{E}_\Omega [\|\mathbf{X} - \hat{\mathbf{X}}\|_*] \leq \left(1 + \frac{r}{R+r+1}\right) \|\mathbf{X} - [\mathbf{X}]_r\|_* \quad \forall r < R \quad (14)$$

where $\|\cdot\|_*$ denotes the nuclear norm and $[\cdot]_r$ is an r -truncated singular-value decomposition of the matrix, which is a best rank- r approximation with respect to every unitarily-invariant norm.

- ▶ \implies We can reduce the storage from $\Theta(p^2)$ to $\Theta(rp)$!

*The algorithm - SketchyCGAL

- ▶ The Augmented Lagrangian of (10) is

$$\mathcal{L}_\beta(\mathbf{X}, \boldsymbol{\lambda}) = \text{Tr}(\mathbf{C}\mathbf{X}) + \langle \boldsymbol{\lambda}, \mathbf{A}\mathbf{X} - \mathbf{b} \rangle + \frac{\beta}{2} \|\mathbf{A}\mathbf{X} - \mathbf{b}\|^2, \quad \nabla_{\mathbf{X}} \mathcal{L}_\beta(\mathbf{X}, \boldsymbol{\lambda}) = \mathbf{C} + \mathbf{A}^T(\boldsymbol{\lambda} + \beta_k(\mathbf{A}\mathbf{X}^k - \mathbf{b}))$$

- ▶ The constraint set of (10) is $\mathcal{X} = \{\mathbf{X} \in \mathbb{R}^{p \times p} : \mathbf{X} \succeq 0, \text{Tr}(\mathbf{X}) = \alpha\}$ and $\text{lmo}_{\mathcal{X}}(\mathbf{Y}) = \alpha v v^T$ where v is the eigenvector corresponding to the minimum eigenvalue of \mathbf{Y} .
- ▶ The algorithm performs linear updates directly on $\mathbf{z}_k := \mathbf{A}\mathbf{X}_k \in \mathbb{R}^n \implies$ the iterates \mathbf{X}_k become **implicit!**

| CGAL | SketchyCGAL (simplified) ⁴ |
|---|--|
| <p>1. Choose $\mathbf{X}^0 = \mathbf{0}_{p \times p} \in \mathcal{X}$, $\boldsymbol{\lambda}^0 = \mathbf{0}_n$, $\beta_0 > 0$, $T > 0$.</p> <p>2. For $k = 0, 1, \dots, T$:</p> $(\xi, v_k) := \text{ApproxMinEvec}(\mathbf{C} + \mathbf{A}^T(\boldsymbol{\lambda}^k + \beta_k(\mathbf{A}\mathbf{X}^k - \mathbf{b})))$ $\mathbf{X}^{k+1} := (1 - \gamma_k)\mathbf{X}^k + \gamma_k(\alpha v_k v_k^T)$ $\boldsymbol{\lambda}^{k+1} := \boldsymbol{\lambda}^k + \omega_k(\mathbf{A}\mathbf{X}^{k+1} - \mathbf{b})$ <p>where $\gamma_k := \frac{2}{k+2}$, and $\beta_k := \frac{\sqrt{k+2}}{\beta_0}$.</p> | <p>1. Choose $\boldsymbol{\lambda}^0 = \mathbf{0}_n$, $\mathbf{z}_0 = \mathbf{0}_n$, $\mathbf{S} = \mathbf{0}_{p \times R}$, $\beta_0 > 0$, $T > 0$, $R > 0$, $\boldsymbol{\Omega} = \text{randn}(p, R)$.</p> <p>2. For $k = 0, 1, \dots, T$:</p> $(\xi, v_k) := \text{ApproxMinEvec}(\mathbf{C} + \mathbf{A}^T(\boldsymbol{\lambda}^k + \beta_k(\mathbf{z}^k - \mathbf{b})))$ $\mathbf{z}^{k+1} := (1 - \gamma_k)\mathbf{z}^k + \gamma_k \mathbf{A}(\alpha v_k v_k^T)$ $\boldsymbol{\lambda}^{k+1} := \boldsymbol{\lambda}^k + w_k(\mathbf{z}^{k+1} - \mathbf{b})$ $\mathbf{S}^{k+1} := (1 - \gamma_k)\mathbf{S}^k + \gamma_k v_k(v_k^T \boldsymbol{\Omega}) \leftarrow \text{update the sketch}$ <p>where $\gamma_k := \frac{2}{k+1}$, and $\beta_k := \frac{\sqrt{k+1}}{\beta_0}$.</p> <p>3. Recover $\hat{\mathbf{X}}_T$ from \mathbf{S}_T using (14)</p> |

*SketchyCGAL: Convergence

- Observations:

- The iterate update procedure of SketchyCGAL is the same as that of CGAL, though \mathbf{X}^k are **implicit**:

$$\begin{aligned}\mathbf{z}^{k+1} &= (1 - \gamma_k)\mathbf{z}^k + \gamma_k \mathbf{A}(\alpha v v^T) \\ \text{by def. of } \mathbf{z}^k &\rightarrow = \mathbf{A}((1 - \gamma_k)\mathbf{X}^k + \gamma_k \alpha v v^T) \\ &= \mathbf{A}\mathbf{X}^{k+1}\end{aligned}$$

- The same computation holds for the sketch updates, where $\mathbf{S}^{k+1} = (1 - \gamma_k)\mathbf{S}^k + \gamma_k \mathbf{v} \mathbf{v}^T \Omega = \mathbf{X}^{k+1} \Omega$.
- the variables in SketchyCGAL track the variables of **some** invocation of CGAL and inherit their behavior.

Theorem [45]

Assume problem (10) satisfies strong duality, and let Ψ^* be its solution set. Then

- The **implicit** iterates converge to the solution set Ψ^* at the same rate as CGAL.
- For each $r < R$, the iterates $\hat{\mathbf{X}}_k$ computed by SketchyCGAL satisfy

$$\lim_{k \rightarrow \infty} \sup \mathbb{E}_\Omega \text{dist}_*(\hat{\mathbf{X}}_k, \Psi^*) \leq \left(1 + \frac{r}{R - r - 1}\right) \max_{\mathbf{Y} \in \Psi^*} \|\mathbf{Y} - [\mathbf{Y}]_r\|_*$$

Here, dist_* is the nuclear-norm distance between a matrix and a set of matrices.

*Example: Convex phase retrieval

Problem formulation

$$f^* := \min_{\mathbf{X} \in \mathbb{C}^{p \times p}} \left\{ \text{Tr}(\mathbf{X}) : \quad \mathcal{A}(\mathbf{X}) = \mathbf{b}, \quad \|\mathbf{X}\|_* \leq \kappa, \quad \mathbf{X} \succeq 0 \right\}. \quad (15)$$

- *This formulation is a convex and semidefinite relaxation of the original, much more difficult Phase Retrieval problem of recovering $\mathbf{x}^\natural \in \mathbb{C}^p$ from the measurements



$$\mathbf{b} \in \mathbb{R}^n, \quad b_i = |\langle \mathbf{a}_i, \mathbf{x}^\natural \rangle|^2 + \omega_i,$$

where $\mathbf{a}_i \in \mathbb{C}^p$ are known measurement vectors, ω_i models noise. Details can be found in [5, 41].

- This type of problem arises, for example, in X-ray crystallography and astronomical imaging.
- Note that the problem is **constrained** to $\mathcal{X} := \{\mathbf{X} \in \mathbb{R}^{p \times p} : \mathbf{X} \succeq 0, \|\mathbf{X}\|_* \leq \kappa\}$, which is **convex** and **compact**.
- \mathcal{X} has an **expensive** prox operator, but an **efficient** lmo.







*Example: Convex Phase Retrieval memory usage

$$f^* := \min_{\mathbf{X} \in \mathbb{C}^{p \times p}} \left\{ \text{Tr}(\mathbf{X}) : \quad \mathcal{A}(\mathbf{X}) = \mathbf{b}, \quad \|\mathbf{X}\|_* \leq \kappa, \quad \mathbf{X} \succeq 0 \right\}.$$



*Example: Max-Cut SDP

$$\max_{\mathbf{X} \in \mathbb{R}^{p \times p}} \left\{ \frac{1}{4} \text{Tr}(\mathbf{L}\mathbf{X}) : \text{diag}(\mathbf{X}) = \mathbf{1}, \mathbf{X} \in \mathcal{S}_+^p, \text{Tr}(\mathbf{X}) = p \right\}$$



Appendix A₁: Generalization of HCGM for $\mathbf{Ax} - \mathbf{b} \in \mathcal{K}$ (self-study)

Quadratic penalty strategy for $\min\{f(\mathbf{x}) : \mathbf{Ax} - \mathbf{b} \in \mathcal{K}, \mathbf{x} \in \mathcal{X}\}$

Define the distance function

$$\text{dist}(\mathbf{y}, \mathcal{K}) := \min_{\mathbf{z} \in \mathcal{K}} \|\mathbf{y} - \mathbf{z}\|.$$

Quadratic penalty takes the form

$$\min_{\mathbf{x} \in \mathbb{R}^p} \left\{ f(\mathbf{x}) + \frac{\beta}{2} \text{dist}^2(\mathbf{Ax} - \mathbf{b}, \mathcal{K}) : \mathbf{x} \in \mathcal{X} \right\}$$

Gradient of $\text{dist}^2(\mathbf{z}, \mathcal{K})$ is

$$\nabla \text{dist}^2(\mathbf{y}, \mathcal{K}) = 2(\mathbf{y} - \text{proj}_{\mathcal{K}}(\mathbf{y})).$$

Hence, HCGM can be generalized by changing lmo step as

$$\hat{\mathbf{x}}^k := \text{lmo}_{\mathcal{X}}(\nabla f(\mathbf{x}^k) + \beta_k \mathbf{A}^T (\mathbf{Ax}^k - \mathbf{b} - \text{proj}_{\mathcal{K}}(\mathbf{Ax}^k - \mathbf{b}))).$$

Same guarantees hold, by replacing $\|\mathbf{Ax} - \mathbf{b}\|$ by $\text{dist}(\mathbf{Ax} - \mathbf{b}, \mathcal{K})$.

Appendix A₂: Generalization of CGAL for $\mathbf{Ax} - \mathbf{b} \in \mathcal{K}$ (self-study)

Augmented Lagrangian for $\min\{f(\mathbf{x}) : \mathbf{Ax} - \mathbf{b} \in \mathcal{K}, \mathbf{x} \in \mathcal{X}\}$

Similarly, CGAL can be extended for $\mathbf{Ax} - \mathbf{b} \in \mathcal{K}$ constraint, by replacing

- ▶ lmo step as

$$\hat{\mathbf{x}}^k := \text{lmo}_{\mathcal{X}} \left(\nabla f(\mathbf{x}^k) + \mathbf{A}^T \lambda^k + \beta_k \mathbf{A}^T (\mathbf{Ax}^k - \mathbf{b} - \text{proj}_{\mathcal{K}}(\mathbf{Ax}^k - \mathbf{b} + \beta_k^{-1} \lambda^k)) \right)$$

- ▶ and dual update step as

$$\lambda^{k+1} := \lambda^k + \omega_k (\mathbf{Ax}^{k+1} - \mathbf{b} + \text{proj}_{\mathcal{K}}(\mathbf{Ax}^{k+1} - \mathbf{b} + \beta_{k+1}^{-1} \lambda^k))$$

Same guarantees hold, by replacing $\|\mathbf{Ax} - \mathbf{b}\|$ by $\text{dist}(\mathbf{Ax} - \mathbf{b}, \mathcal{K})$.

References I

- [1] Sanjeev Arora, Satish Rao, and Umesh Vazirani.
Expander flows, geometric embeddings and graph partitioning.
Journal of the ACM (JACM), 56(2):5, 2009.
- [2] H.H. Bauschke and P. Combettes.
Convex analysis and monotone operators theory in Hilbert spaces.
Springer-Verlag, 2011.
- [3] Dimitri P Bertsekas.
Necessary and sufficient conditions for a penalty method to be exact.
Mathematical programming, 9(1):87–99, 1975.
- [4] Dimitri P Bertsekas.
On penalty and multiplier methods for constrained minimization.
SIAM Journal on Control and Optimization, 14(2):216–235, 1976.
- [5] Emmanuel J Candes, T. Strohmer, and V. Voroninski.
Phaselift: Exact and stable signal recovery from magnitude measurements via convex programming.
IEEE Trans. Signal Processing, 60(5):2422–2432, 2012.

References II

- [6] Coralia Cartis, Nicholas IM Gould, and Ph L Toint.
Complexity bounds for second-order optimality in unconstrained optimization.
Journal of Complexity, 28(1):93–108, 2012.
- [7] A. Chambolle and T. Pock.
A first-order primal-dual algorithm for convex problems with applications to imaging.
Journal of Mathematical Imaging and Vision, 40(1):120–145, 2011.
- [8] Vaggos Chatziafratis, Rad Niazadeh, and Moses Charikar.
Hierarchical clustering with structural constraints.
arXiv preprint arXiv:1805.09476, 2018.
- [9] G. Chen and M. Teboulle.
A proximal-based decomposition method for convex minimization problems.
Math. Program., 64:81–101, 1994.
- [10] P. L. Combettes and V. R. Wajs.
Signal recovery by proximal forward-backward splitting.
Multiscale Model. Simul., 4:1168–1200, 2005.

References III

- [11] D. Davis.
Convergence rate analysis of the forward-Douglas-Rachford splitting scheme.
UCLA CAM report 14-73, 2014.
- [12] D. Davis and W. Yin.
Faster convergence rates of relaxed Peaceman-Rachford and ADMM under regularity assumptions.
UCLA CAM report 14-58, 2014.
- [13] D. Davis and W. Yin.
A three-operator splitting scheme and its optimization applications.
Tech. Report., 2015.
- [14] J. Eckstein and D. Bertsekas.
On the Douglas - Rachford splitting method and the proximal point algorithm for maximal monotone operators.
Math. Program., 55:293–318, 1992.
- [15] J. E. Esser.
Primal-dual algorithm for convex models and applications to image restoration, registration and nonlocal inpainting.
Phd. thesis, University of California, Los Angeles, Los Angeles, USA, 2010.

References IV

- [16] D. Gabay and B. Mercier.
A dual algorithm for the solution of nonlinear variational problems via finite element approximation.
Computers & Mathematics with Applications, 2(1):17 – 40, 1976.
- [17] Saeed Ghadimi and Guanghui Lan.
Accelerated gradient methods for nonconvex nonlinear and stochastic programming.
Math. Program., 156(1-2):59–99, 2016.
- [18] T. Goldstein, E. Esser, and R. Baraniuk.
Adaptive Primal-Dual Hybrid Gradient Methods for Saddle Point Problems.
Tech. Report., <http://arxiv.org/pdf/1305.0546v1.pdf>:1–26, 2013.
- [19] T. Goldstein, B. ODonoghue, and S. Setzer.
Fast Alternating Direction Optimization Methods.
SIAM J. Imaging Sci., 7(3):1588–1623, 2012.
- [20] B. He and X. Yuan.
Convergence analysis of primal-dual algorithms for saddle-point problem: from contraction perspective.
SIAM J. Imaging Sciences, 5:119–149, 2012.

References V

- [21] Bingsheng He and Xiaoming Yuan.
On the acceleration of augmented lagrangian method for linearly constrained optimization.
2010.
- [22] B.S. He and X.M. Yuan.
On the $O(1/n)$ convergence rate of the Douglas-Rachford alternating direction method.
SIAM J. Numer. Anal., 50:700–709, 2012.
- [23] Magnus R Hestenes.
Multiplier and gradient methods.
Journal of optimization theory and applications, 4(5):303–320, 1969.
- [24] Dustin G Mixon, Soledad Villar, and Rachel Ward.
Clustering subgaussian mixtures by semidefinite programming.
arXiv preprint arXiv:1602.06612, 2016.
- [25] Aryan Mokhtari, Hamed Hassani, and Amin Karbasi.
Stochastic conditional gradient methods: From convex minimization to submodular maximization.
arXiv preprint arXiv:1804.09554, 2018.
- [26] I. Necoara and J.A.K. Suykens.
Interior-point lagrangian decomposition method for separable convex optimization.
J. Optim. Theory and Appl., 143(3):567–588, 2009.

References VI

- [27] Y. Ouyang, Y. Chen, G. LanG. Lan., and E. JR. Pasiliao.
An accelerated linearized alternating direction method of multiplier.
Tech, 2014.
- [28] Yuyuan Ouyang and Yangyang Xu.
Lower complexity bounds of first-order methods for convex-concave bilinear saddle-point problems.
arXiv preprint arXiv:1808.02901, 2018.
- [29] Jiming Peng and Yu Wei.
Approximating k-means-type clustering via semidefinite programming.
SIAM journal on optimization, 18(1):186–205, 2007.
- [30] Michael JD Powell.
A method for nonlinear constraints in minimization problems.
Optimization, pages 283–298, 1969.
- [31] Shoham Sabach and Marc Teboulle.
Lagrangian methods for composite optimization.
In *Handbook of Numerical Analysis*, volume 20, pages 401–436. Elsevier, 2019.

References VII

- [32] Mehmet Fatih Sahin, Ahmet Alacaoglu, Fabian Latorre, Volkan Cevher, et al.
An inexact augmented lagrangian framework for nonconvex optimization with nonlinear constraints.
In Advances in Neural Information Processing Systems, pages 13965–13977, 2019.
- [33] Defeng Sun, Kim-Chuan Toh, Yancheng Yuan, and Xin-Yuan Zhao.
Sdpnal+: A matlab software for semidefinite programming with bound constraints (version 1.0).
Optimization Methods and Software, 35(1):87–115, 2020.
- [34] Q. Tran-Dinh, I. Necoara, C. Savorgnan, and M. Diehl.
Asymptotic pseudotrajectories and chain recurrent flows, with applications.
SIAM J. Optim., 8(1):141–176, 1996.
- [35] Quoc Tran-Dinh, Ahmet Alacaoglu, Olivier Fercoq, and Volkan Cevher.
An adaptive primal-dual framework for nonsmooth convex minimization.
arXiv preprint arXiv:1808.04648, 2018.
- [36] Quoc Tran-Dinh and Volkan Cevher.
Constrained convex minimization via model-based excessive gap.
In Proceedings of the 27th International Conference on Neural Information Processing Systems - Volume 1, NIPS'14, 2014.

References VIII

- [37] P. Tseng.
Applications of splitting algorithm to decomposition in convex programming and variational inequalities.
SIAM J. Control Optim., 29:119–138, 1991.
- [38] E. Wei, A. Ozdaglar, and A.Jadbabaie.
A Distributed Newton Method for Network Utility Maximization.
<http://web.mit.edu/asuman/www/publications.htm>, 2011.
- [39] Yangyang Xu.
Accelerated first-order primal-dual proximal methods for linearly constrained composite convex programming.
SIAM Journal on Optimization, 27(3):1459–1484, 2017.
- [40] Alp Yurtsever, Olivier Fercoq, and Volkan Cevher.
A conditional gradient-based augmented lagrangian framework.
Technical report, 2018.
- [41] Alp Yurtsever, Ya-Ping Hsieh, and Volkan Cevher.
Scalable convex methods for phase retrieval.
In *6th IEEE Intl. Workshop on Computational Advances in Multi-Sensor Adaptive Processing*, 2015.

References IX

- [42] Alp Yurtsever, Fercoq Olivier, Locatello Francesco, and Volkan Cevher.
A conditional gradient framework for composite convex minimization with applications to semidefinite programming.
In *Proceedings of the 35th International Conference on International Conference on Machine Learning - Volume 28*, ICML'18, 2018.
- [43] Alp Yurtsever, Suvrit Sra, and Volkan Cevher.
Conditional gradient methods via stochastic path-integrated differential estimator.
In *International Conference on Machine Learning*, pages 7282–7291. PMLR, 2019.
- [44] Alp Yurtsever, Quoc Tran-Dinh, and Volkan Cevher.
A universal primal-dual convex optimization framework.
In *Proceedings of the 28th International Conference on Neural Information Processing Systems - Volume 2*, NIPS'15, 2015.
- [45] Alp Yurtsever, Joel A Tropp, Olivier Fercoq, Madeleine Udell, and Volkan Cevher.
Scalable semidefinite programming.
arXiv preprint arXiv:1912.02949, 2019.